

# Multi-Layer Perceptrons and Deep Learning

Christos Dimitrakakis

October 23, 2024

# Outline

## Features and layers

## Algorithms

- Random projection

- Back propagation

- Derivatives

- Cost functions

- Stochastic gradient descent in practice

## Python libraries

- sklearn

- PyTorch

- TensorFlow

## Features and layers

### Algorithms

Random projection

Back propagation

Derivatives

Cost functions

Stochastic gradient descent in practice


### Python libraries

sklearn

PyTorch

TensorFlow

# Perceptron vs linear regression



- Network output

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

- Chain rule

$$\nabla_{\beta} L = \nabla_{\hat{y}} L \nabla_{\beta} \hat{y}$$

- Network gradient

$$\nabla_{\beta} \hat{y} = (x_1, x_2)$$

## Cost functions

The only difference are the cost functions

- Perceptron

$$L = -\mathbb{I}\{y \neq \hat{y}\} \hat{y}$$

with

$$\nabla L = -\mathbb{I}\{y \neq \hat{y}\} yx$$

- Linear regression

$$L = (\hat{y} - y)^2,$$

with

$$\nabla_{\hat{y}} L = 2(\hat{y} - y).$$

# Layering and features

## Fixed layers

- ▶ Input to layer  $x \in R^n$
- ▶ Output from layer  $\hat{y} \in R^m$ .

## Intermediate layers


- ▶ Linear layer
- ▶ Non-linear **activation** function.

## Linear layers types

- ▶ Dense
- ▶ Sparse
- ▶ Convolutional

## Activation function

- ▶ Sigmoid
- ▶ Softmax



Input layer

Linear layer

Sigmoid activation

Linear layer

Softmax activation

# Linear layers

Example: Linear regression with  $n$  inputs,  $m$  outputs.

- ▶ Input: Features  $\mathbf{x} \in \mathbb{R}^n$
- ▶ Dense linear layer with  $\mathbf{B} \in \mathbb{R}^{m \times n}$
- ▶ Output:  $\hat{\mathbf{y}} \in \mathbb{R}^m$

## Dense linear layer

- ▶ Parameters  $\mathbf{B} = \begin{pmatrix} \beta_1 \\ \vdots \\ \beta_m \end{pmatrix}$ ,
- ▶  $\beta_i = [\beta_{i,1}, \dots, \beta_{i,n}]$ ,  $\beta_i$  connects the  $i$ -th output  $y_i$  to the features  $\mathbf{x}$ :  

$$y_i = \beta_i \mathbf{x}$$
- ▶ In compact form:

$$\mathbf{y} = \mathbf{B}\mathbf{x}$$

# ReLU layers

- Typically used in the hidden layers of neural networks

$$f(x) = \max(0, x)$$

## Derivative

$$df/dxf(x) = \mathbb{I}\{x > 0\}$$

# Sigmoid activation

## Example: Logistic regression


- ▶ Input  $\mathbf{x} \in \mathbb{R}^n$
- ▶ Intermediate output:  $z \in \mathbb{R}$ ,

$$z = \sum_{i=1}^n \beta_i x_i.$$

- ▶ Output: sigmoid activation  
 $\hat{y} \in [0, 1]$ .

$$f(z) = 1/[1 + \exp(-z)].$$

Now we can interpret  $\hat{y} = P_{\beta}(y = 1|x)$ .



Input layer

Linear layer

Sigmoid layer

## Loss function: negative log likelihood

$$\ell(\hat{y}, y) = -[\mathbb{I}\{y = 1\} \ln(\hat{y}) + \mathbb{I}\{y = -1\} \ln(1 - \hat{y})]$$



# Softmax layer

Example: Multivariate logistic regression with  $m$  classes.

- ▶ Input: **Features**  $\mathbf{x} \in \mathbb{R}^n$
- ▶ Fully-connected **linear** activation layer

$$\mathbf{z} = \mathbf{B}\mathbf{x}, \quad \mathbf{B} \in \mathbb{R}^{m \times n}.$$


- ▶ **Softmax** output

$$\hat{y}_i = \frac{\exp(z_i)}{\sum_{j=1}^m \exp(z_j)}$$

We can also interpret this as

$$\hat{y}_i \triangleq \mathbb{P}(y = i \mid \mathbf{x})$$

with usual loss  $\ell(\hat{\mathbf{y}}, \mathbf{y}) = -\ln \hat{y}_y$



Input layer

Linear layer

Softmax layer

## Features and layers

## Algorithms

- Random projection

- Back propagation

- Derivatives

- Cost functions

- Stochastic gradient descent in practice

## Python libraries

- sklearn

- PyTorch

- TensorFlow

# Random projections

- ▶ Features  $x$
- ▶ Hidden layer activation  $z$
- ▶ Output  $y$

## Hidden layer: Random projection

Here we project the input into a high-dimensional space

$$z_i = \text{sgn}(\beta_i^\top x) = y_i$$

where  $B = [\beta_i]_{i=1}^m$ ,  $\beta_{i,j} \sim \text{Normal}(0, 1)$

## The reason for random projections

- ▶ The high dimension makes it easier to learn.
- ▶ The randomness ensures we are not learning something spurious.

# Background on back-propagation

## Gradient descent algorithm

- ▶ We need to minimise the expected value  $\mathbb{E}_{\beta}[L]$  of the loss function  $L$
- ▶ Since we cannot calculate  $\mathbb{E}_{\beta}[L]$ , we use:

$$\nabla_{\beta} \mathbb{E}_{\beta}[L] \approx \frac{1}{T} \sum_{t=1}^T \nabla_{\beta} \ell(x_t, y_t, \beta).$$

- ▶ We can then update our parameters to reduce the **empirical loss**

$$\beta_{t+1} = \beta_t - \alpha_t \nabla_{\beta} \ell(x_t, y_t, \beta).$$

## The problem

- ▶ However  $\ell$  is a complex function of  $\beta$ .
- ▶ How can we obtain  $\nabla_{\beta} \ell$ ?

## The solution


- ▶ Use the chain rule to "backpropagate" errors.

# The chain rule of differentiation



[1673] Leibniz

## Chain rule applied to the perceptron



[1976] Rosenblat

# Chain rule for deep neural networks



[1982] Werbos

# Backpropagation given a name


1986: Learning representations by back-propagating errors.



Rumelhart




Hinton



Williams




# Elementary back-propagation: linear regression



- $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$ , chain rule:  $\nabla_\beta \ell = \nabla_\beta f \nabla_{\hat{y}} \ell$
- **Forward**: follow the arrows to calculate **variables**

$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$


# Elementary back-propagation: linear regression



- $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$ , chain rule:  $\nabla_\beta \ell = \nabla_\beta f \nabla_{\hat{y}} \ell$
- **Forward**: follow the arrows to calculate **variables**

$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$


# Elementary back-propagation: linear regression



- $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$ , chain rule:  $\nabla_{\beta} \ell = \nabla_{\beta} f \nabla_{\hat{y}} \ell$
- **Forward**: follow the arrows to calculate **variables**

$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$

# Elementary back-propagation: linear regression



►  $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$ , chain rule:  $\nabla_{\beta} \ell = \nabla_{\beta} f \nabla_{\hat{y}} \ell$

► **Forward:** follow the arrows to calculate **variables**


$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$

► **Backward:** return to calculate the **gradients**

$$\nabla_{\beta} \ell(\hat{y}, y) = \nabla_{\beta} f(\beta, x) \times \nabla_{\hat{y}} \ell(\hat{y}, y) \quad (1)$$

$$= \nabla_{\beta} f(\beta, x) \times 2[\hat{y} - y] \quad (2)$$

# Elementary back-propagation: linear regression



►  $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$ , chain rule:  $\nabla_{\beta} \ell = \nabla_{\beta} f \nabla_{\hat{y}} \ell$

► **Forward:** follow the arrows to calculate **variables**


$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$

► **Backward:** return to calculate the **gradients**

$$\nabla_{\beta} \ell(\hat{y}, y) = \nabla_{\beta} f(\beta, x) \times \nabla_{\hat{y}} \ell(\hat{y}, y) \quad (1)$$

$$= \nabla_{\beta} f(\beta, x) \times 2[\hat{y} - y] \quad (2)$$

# Elementary back-propagation: linear regression



►  $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$ , chain rule:  $\nabla_{\beta} \ell = \nabla_{\beta} f \nabla_{\hat{y}} \ell$

► **Forward:** follow the arrows to calculate **variables**

$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$

► **Backward:** return to calculate the **gradients**

$$\nabla_{\beta} \ell(\hat{y}, y) = \nabla_{\beta} f(\beta, x) \times \nabla_{\hat{y}} \ell(\hat{y}, y) \quad (1)$$

$$= \nabla_{\beta} f(\beta, x) \times 2[\hat{y} - y] \quad (2)$$

► Update:

$$\beta_{t+1} = \beta_t - \alpha_t \times \nabla_{\beta} \ell(\hat{y}_t, y_t)$$


## Gradient descent with *back-propagation*

- ▶ Dataset  $D$ , cost function  $L = \sum_t \ell_t$
- ▶ Parameters  $B_1, \dots, B_k$  with  $k$  layers
- ▶ Intermediate variables:  $z_j = h_j(z_{j-1}, B_j)$ ,  $z_0 = x$ ,  $z_k = \hat{y}$ .

## Gradient descent with *back-propagation*

- ▶ Dataset  $D$ , cost function  $L = \sum_t \ell_t$
- ▶ Parameters  $B_1, \dots, B_k$  with  $k$  layers
- ▶ Intermediate variables:  $z_j = h_j(z_{j-1}, B_j)$ ,  $z_0 = x$ ,  $z_k = \hat{y}$ .

### Dependency graph






# Gradient descent with *back-propagation*

- ▶ Dataset  $D$ , cost function  $L = \sum_t \ell_t$
- ▶ Parameters  $B_1, \dots, B_k$  with  $k$  layers
- ▶ Intermediate variables:  $z_j = h_j(z_{j-1}, B_j)$ ,  $z_0 = x$ ,  $z_k = \hat{y}$ .

## Dependency graph



## Backpropagation with steepest stochastic gradient descent

- ▶ Forward step: For  $j = 1, \dots, k$ , calculate  $z_j = h_j(k)$  and  $\ell(\hat{y}, y)$
- ▶ Backward step: Calculate  $\nabla_{\hat{y}} \ell$  and  $d_j \triangleq \nabla_{B_j} \ell = \nabla_{B_j} z_j d_{j+1}$  for  $j = k \dots, 1$
- ▶ Apply gradient:  $B_j \leftarrow B_j - \alpha d_j$ .

# Other algorithms and gradients

## Natural gradient

Defined for probabilistic models

## ADAM

Exponential moving average of gradient and square gradients

## BFGS: Broyden–Fletcher–Goldfarb–Shanno algorithm

Newton-like method

# Linear layer

## Definition

This is a linear combination of inputs  $x \in \mathbb{R}^n$  and parameter matrix  $B \in \mathbb{R}^{m \times n}$

$$\text{where } B = \begin{bmatrix} \beta_1 \\ \vdots \\ \beta_i \\ \vdots \\ \beta_m \end{bmatrix} = \begin{bmatrix} \beta_{1,1} & \cdots & \beta_{1,j} & \cdots & \beta_{1,m} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \beta_{i,1} & \cdots & \beta_{i,j} & \cdots & \beta_{i,m} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \beta_{n,1} & \cdots & \beta_{n,j} & \cdots & \beta_{n,m} \end{bmatrix}$$

$$f(B, x) = Bx \quad f_i(B, x) = \beta_i \cdot x = \sum_{j=1}^n \beta_{i,j} x_j,$$

## Gradient

Each partial derivative is simple:

$$\frac{\partial}{\partial \beta_{i,j}} f_k(B, x) = \sum_{k=1}^n \frac{\partial}{\partial \beta_{i,j}} \beta_{i,k} x_k = x_j$$

## Sigmoid layer

- ▶ This layer is used for **binary classification**.
- ▶ It is used in the **logistic regression** model to obtain label probabilities.

$$f(z) = 1/(1 + \exp(-z))$$

## Derivative

So let us ignore the other inputs for simplicity:

$$\frac{d}{dz} f(z) = \exp(-z)/[1 + \exp(-z)]^2$$

# Softmax layer

- This layer is used for **multi-class classification**

$$y_i(z) = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

## Derivative

$$\frac{\partial}{\partial z_i} y_i(z) = \frac{e^{z_i} e^{\sum_{j \neq i} z_j}}{\left(\sum_j e^{z_j}\right)^2}$$

$$\frac{\partial}{\partial z_i} y_k(z) = \frac{e^{z_i + z_k}}{\left(\sum_j e^{z_j}\right)^2}$$

# Classification cost functions

## Classification error

If  $z$  is the output for each class then

$$\ell(z, y) = \mathbb{I}\{y \notin \arg \max(z)\}$$

This is not differentiable.

## Error margin

If  $z$  is the positive class output then

$$\ell(z, y) = -\mathbb{I}\{zy < 0\} zy$$


Used in the perceptron.

## Negative log likelihood

If  $z$  are label probabilities, then

$$\ell(z, y) = -\ln z_y.$$

Used in logistic regression.



## Hinge loss

If  $z$  are the output for each class

$$\ell(z, y) = 1 - z_y$$

Used in large margin classifiers.


# Regression cost functions

## L2 loss (Squared error)

If  $z$  is a prediction for  $y$  then

$$\ell(z, y) = (y - z)^2$$

This is equivalent to negative log likelihood under Gaussianity. Used in linear regression.



## L1 loss

If  $z$  is a prediction for  $y$  then

$$\ell(z, y) = |y - z|$$

Used in LASSO regression.

## Huber loss

If  $z$  is a prediction, then

$$\ell(z, y) = \begin{cases} \frac{1}{2}(z - y)^2 & |z - y| \leq \delta \\ \delta(|z - y| - \frac{1}{2}\delta) & \text{otherwise.} \end{cases} \quad (3)$$

Mixes L1 and L2 losses.

# Gradient descent in practice

## The ideal gradient descent algorithm:

If we could calculate  $\nabla_{\beta} \mathbb{E}_{\beta}[L]$ , we could do:

$$\beta_{n+1} = \beta_n - \alpha_n \nabla_{\beta} \mathbb{E}_{\beta}[L]$$

for a suitable  $\alpha_n$  schedule.

## Gradient descent on the empirical error

Since we only have the data, we can try to minimise the empirical loss  $\frac{1}{T} \sum_{t=1}^T \ell(x_t, y_t, \beta)$  through gradient descent

$$\beta_{n+1} = \beta_n - \alpha_n \frac{1}{T} \sum_{t=1}^T \nabla_{\beta} \ell(x_t, y_t, \beta)$$

This is also called **batch** gradient descent.



# Stochastic gradient descent

## Gradient descent on one example:

We don't have to wait calculate  $\nabla_{\beta} \ell(x_t, y_t, \beta)$  for all  $t$  before applying the update. We can do it at every example:

$$\beta_{n+1} = \beta_n - \alpha_n \nabla_{\beta} \ell(x_{[n]_T}, y_{[n]_T}, \beta).$$

Here  $[n]_T$  is  $1 + n$  modulo  $T$  to ensure  $n \in \{1, \dots, T\}$ .

## Minibatch gradient descent

However, it is a bit better to look at  $K$  examples at a time before we change the parameters. This is called a **minibatch**

$$\beta_{n+1} = \beta_n - \alpha_n \frac{1}{K} \sum_{k=nK}^{(n+1)K-1} \nabla_{\beta} \ell(x_{[k]_T}, y_{[k]_T}, \beta)$$

This also helps with parallelisation, since we can compute  $\ell, \nabla_{\beta} \ell$  in parallel for each example.

## Features and layers

## Algorithms

- Random projection

- Back propagation

- Derivatives

- Cost functions

- Stochastic gradient descent in practice

## Python libraries

- sklearn

- PyTorch

- TensorFlow

# sklearn neural networks

## Classification

Uses the **cross entropy** cost

```
from sklearn.neural_network import MLPClassifier
clf = MLPClassifier(hidden_layer_sizes=(5, 2))
clf.fit(X, y)
clf.predict(X_test)
```

- ▶ Main condition is layer sizes.

## Regression

```
from sklearn.neural_network import MLPRegressor
model = MLPRegressor(hidden_layer_sizes=(5, 2))
```

# PyTorch

## Data set-up

```
X_train = torch.tensor(X_train, dtype=torch.float32)
train_dataset = TensorDataset(X_train, y_train)
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
```

# PyTorch: Manual training

## Network setup

```
fc1 = nn.Linear(input_size, hidden_size) # Input to hidden layer
fc2 = nn.Linear(hidden_size, output_size) # Hidden layer to output
sigmoid = nn.Sigmoid() # some activation function
criterion = nn.BCELoss() #what loss to minimise
optimizer = optim.SGD(model.parameters(), lr=0.001) # how to minimise
```

## Training

```
# Manual forward pass.
z1 = fc1(inputs) # hidden layer 1
a1 = sigmoid(z1) # Apply activation for hidden
z2 = fc2(a1) # Linear combination in output layer
outputs = sigmoid(z2) # Output layer activation
loss = criterion(outputs, labels) # Specify loss
loss.backward() # Backward pass
optimizer.step() # Update weights
```

# TensorFlow

This is another library, no need to use this for this course