

Multi-Layer Perceptrons and Deep Learning

Christos Dimitrakakis

October 22, 2024

Outline

Features and layers

Algorithms

- Random projection

- Back propagation

- Derivatives

- Cost functions

Python libraries

- sklearn

- PyTorch

- TensorFlow

Features and layers

Algorithms

Random projection

Back propagation

Derivatives

Cost functions


Python libraries

sklearn

PyTorch

TensorFlow

Perceptron vs linear regression



- Network output

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

- Chain rule

$$\nabla_{\beta} L = \nabla_{\hat{y}} L \nabla_{\beta} \hat{y}$$

- Network gradient

$$\nabla_{\beta} \hat{y} = (x_1, x_2)$$

Cost functions

The only difference are the cost functions

- Perceptron

$$L = -\mathbb{I}\{y \neq \hat{y}\} \hat{y}$$

with

$$\nabla L = -\mathbb{I}\{y \neq \hat{y}\} y x$$

- Linear regression

$$L = (\hat{y} - y)^2,$$

with

$$\nabla_{\hat{y}} L = 2(\hat{y} - y).$$

Layering and features

Fixed layers

- ▶ Input to layer $x \in R^n$
- ▶ Output from layer $\hat{y} \in R^m$.

Intermediate layers


- ▶ Linear layer
- ▶ Non-linear **activation** function.

Linear layers types

- ▶ Dense
- ▶ Sparse
- ▶ Convolutional

Activation function

- ▶ Sigmoid
- ▶ Softmax



Input layer

Linear layer

Sigmoid activation

Linear layer

Softmax activation

Linear layers

Example: Linear regression with n inputs, m outputs.

- ▶ Input: Features $\mathbf{x} \in \mathbb{R}^n$
- ▶ Dense linear layer with $\mathbf{B} \in \mathbb{R}^{m \times n}$
- ▶ Output: $\hat{\mathbf{y}} \in \mathbb{R}^m$

Dense linear layer

- ▶ Parameters $\mathbf{B} = \begin{pmatrix} \beta_1 \\ \vdots \\ \beta_m \end{pmatrix}$,
- ▶ $\beta_i = [\beta_{i,1}, \dots, \beta_{i,n}]$, β_i connects the i -th output y_i to the features \mathbf{x} :

$$y_i = \beta_i \mathbf{x}$$

- ▶ In compact form:

$$\mathbf{y} = \mathbf{B}\mathbf{x}$$

ReLU layers

- ▶ Typically used in the hidden layers of neural networks

$$f(x) = \max(0, x)$$

Derivative

$$df/dxf(x) = \mathbb{I}\{x > 0\}$$

Sigmoid activation

Example: Logistic regression


- ▶ Input $\mathbf{x} \in \mathbb{R}^n$
- ▶ Intermediate output: $z \in \mathbb{R}$,

$$z = \sum_{i=1}^n \beta_i x_i.$$

- ▶ Output: sigmoid activation
 $\hat{y} \in [0, 1]$.

$$f(z) = 1/[1 + \exp(-z)].$$

Now we can interpret $\hat{y} = P_{\beta}(y = 1|x)$.



Input layer

Linear layer

Sigmoid layer

Loss function: negative log likelihood

$$\ell(\hat{y}, y) = -[\mathbb{I}\{y = 1\} \ln(\hat{y}) + \mathbb{I}\{y = -1\} \ln(1 - \hat{y})]$$

Softmax layer

Example: Multivariate logistic regression with m classes.

- ▶ Input: **Features** $\mathbf{x} \in \mathbb{R}^n$
- ▶ Fully-connected **linear** activation layer

$$\mathbf{z} = \mathbf{B}\mathbf{x}, \quad \mathbf{B} \in \mathbb{R}^{m \times n}.$$


- ▶ **Softmax** output

$$\hat{y}_i = \frac{\exp(z_i)}{\sum_{j=1}^m \exp(z_j)}$$

We can also interpret this as

$$\hat{y}_i \triangleq \mathbb{P}(y = i \mid \mathbf{x})$$

with usual loss $\ell(\hat{y}, y) = -\ln \hat{y}_y$



Input layer

Linear layer

Softmax layer

Features and layers

Algorithms

- Random projection

- Back propagation

- Derivatives

- Cost functions

Python libraries

- sklearn

- PyTorch

- TensorFlow

Random projections

- ▶ Features x
- ▶ Hidden layer activation z
- ▶ Output y

Hidden layer: Random projection

Here we project the input into a high-dimensional space

$$z_i = \text{sgn}(\beta_i^\top x) = y_i$$

where $B = [\beta_i]_{i=1}^m$, $\beta_{i,j} \sim \text{Normal}(0, 1)$

The reason for random projections

- ▶ The high dimension makes it easier to learn.
- ▶ The randomness ensures we are not learning something spurious.

Background on back-propagation

The problem

- ▶ We need to minimise a loss function L
- ▶ We need to calculate

$$\nabla_{\beta} \mathbb{E}_{\beta}[L] \approx \frac{1}{T} \sum_{t=1}^T \nabla_{\beta} \ell(x_t, y_t, \beta).$$


- ▶ However $\ell(x_t, y_t, \beta)$ is a complex non-linear function of β .
- ▶ We need many steps to calculate ℓ . How can we then do it?

The chain rule of differentiation



[1673] Leibniz

Chain rule applied to the perceptron



[1976] Rosenblatt

Chain rule for deep neural networks



[1982] Werbos

Backpropagation given a name


1986: Learning representations by back-propagating errors.



Rumelhart




Hinton



Williams


Elementary back-propagation: linear regression



- $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$, chain rule: $\nabla_\beta \ell = \nabla_\beta f \nabla_{\hat{y}} \ell$
- **Forward**: follow the arrows to calculate **variables**

$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$


Elementary back-propagation: linear regression



- $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$, chain rule: $\nabla_\beta \ell = \nabla_\beta f \nabla_{\hat{y}} \ell$
- **Forward**: follow the arrows to calculate **variables**

$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$


Elementary back-propagation: linear regression



- $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$, chain rule: $\nabla_{\beta} \ell = \nabla_{\beta} f \nabla_{\hat{y}} \ell$
- **Forward**: follow the arrows to calculate **variables**

$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$

Elementary back-propagation: linear regression



- $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$, chain rule: $\nabla_{\beta} \ell = \nabla_{\beta} f \nabla_{\hat{y}} \ell$
- **Forward**: follow the arrows to calculate **variables**


$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$

- **Backward**: return to calculate the **gradients**

$$\nabla_{\beta} \ell(\hat{y}, y) = \nabla_{\beta} f(\beta, x) \times \nabla_{\hat{y}} \ell(\hat{y}, y) \quad (1)$$

$$= \nabla_{\beta} f(\beta, x) \times 2[\hat{y} - y] \quad (2)$$

Elementary back-propagation: linear regression



- $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$, chain rule: $\nabla_{\beta} \ell = \nabla_{\beta} f \nabla_{\hat{y}} \ell$
- **Forward**: follow the arrows to calculate **variables**


$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$

- **Backward**: return to calculate the **gradients**

$$\nabla_{\beta} \ell(\hat{y}, y) = \nabla_{\beta} f(\beta, x) \times \nabla_{\hat{y}} \ell(\hat{y}, y) \quad (1)$$

$$= \nabla_{\beta} f(\beta, x) \times 2[\hat{y} - y] \quad (2)$$

Elementary back-propagation: linear regression



► $f : X \rightarrow Y, \ell : Y \times Y \rightarrow \mathbb{R}$, chain rule: $\nabla_{\beta} \ell = \nabla_{\beta} f \nabla_{\hat{y}} \ell$

► **Forward:** follow the arrows to calculate **variables**

$$\hat{y} \triangleq f(\beta, x) = \sum_{i=1}^n \beta_i x_i, \quad \ell(\hat{y}, y) = (\hat{y} - y)^2$$

► **Backward:** return to calculate the **gradients**

$$\nabla_{\beta} \ell(\hat{y}, y) = \nabla_{\beta} f(\beta, x) \times \nabla_{\hat{y}} \ell(\hat{y}, y) \quad (1)$$

$$= \nabla_{\beta} f(\beta, x) \times 2[\hat{y} - y] \quad (2)$$

► Update:

$$\beta_{t+1} = \beta_t - \alpha_t \times \nabla_{\beta} \ell(\hat{y}_t, y_t)$$


Gradient descent with *back-propagation*

- ▶ Dataset D , cost function $L = \sum_t \ell_t$
- ▶ Parameters B_1, \dots, B_k with k layers
- ▶ Intermediate variables: $z_j = h_j(z_{j-1}, B_j)$, $z_0 = x$, $z_k = \hat{y}$.

Gradient descent with *back-propagation*

- ▶ Dataset D , cost function $L = \sum_t \ell_t$
- ▶ Parameters B_1, \dots, B_k with k layers
- ▶ Intermediate variables: $z_j = h_j(z_{j-1}, B_j)$, $z_0 = x$, $z_k = \hat{y}$.


Dependency graph



Gradient descent with *back-propagation*

- ▶ Dataset D , cost function $L = \sum_t \ell_t$
- ▶ Parameters B_1, \dots, B_k with k layers
- ▶ Intermediate variables: $z_j = h_j(z_{j-1}, B_j)$, $z_0 = x$, $z_k = \hat{y}$.

Dependency graph



Backpropagation with steepest stochastic gradient descent

- ▶ Forward step: For $j = 1, \dots, k$, calculate $z_j = h_j(k)$ and $\ell(\hat{y}, y)$
- ▶ Backward step: Calculate $\nabla_{\hat{y}} \ell$ and $d_j \triangleq \nabla_{B_j} \ell = \nabla_{B_j} z_j d_{j+1}$ for $j = k \dots, 1$
- ▶ Apply gradient: $B_j \leftarrow B_j - \alpha d_j$.

Other algorithms and gradients

Natural gradient

Defined for probabilistic models

ADAM

Exponential moving average of gradient and square gradients

BFGS: Broyden–Fletcher–Goldfarb–Shanno algorithm

Newton-like method

Linear layer

Definition

This is a linear combination of inputs $x \in \mathbb{R}^n$ and parameter matrix $B \in \mathbb{R}^{m \times n}$

$$\text{where } B = \begin{bmatrix} \beta_1 \\ \vdots \\ \beta_i \\ \vdots \\ \beta_m \end{bmatrix} = \begin{bmatrix} \beta_{1,1} & \cdots & \beta_{1,j} & \cdots & \beta_{1,m} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \beta_{i,1} & \cdots & \beta_{i,j} & \cdots & \beta_{i,m} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ \beta_{n,1} & \cdots & \beta_{n,j} & \cdots & \beta_{n,m} \end{bmatrix}$$

$$f(B, x) = Bx \quad f_i(B, x) = \beta_i \cdot x = \sum_{j=1}^n \beta_{i,j} x_j,$$

Gradient

Each partial derivative is simple:

$$\frac{\partial}{\partial \beta_{i,j}} f_k(B, x) = \sum_{k=1}^n \frac{\partial}{\partial \beta_{i,j}} \beta_{i,k} x_k = x_j$$

Sigmoid layer

- ▶ This layer is used for **binary classification**.
- ▶ It is used in the **logistic regression** model to obtain label probabilities.

$$f(z) = 1/(1 + \exp(-z))$$

Derivative

So let us ignore the other inputs for simplicity:

$$\frac{d}{dz} f(z) = \exp(-z)/[1 + \exp(-z)]^2$$

Softmax layer

- This layer is used for **multi-class classification**

$$y_i(z) = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

Derivative

$$\frac{\partial}{\partial z_i} y_i(z) = \frac{e^{z_i} e^{\sum_{j \neq i} z_j}}{\left(\sum_j e^{z_j}\right)^2}$$

$$\frac{\partial}{\partial z_i} y_k(z) = \frac{e^{z_i + z_k}}{\left(\sum_j e^{z_j}\right)^2}$$

Classification cost functions

Error margin

If z is a confidence level for the positive class then

$$\ell(z, y) = -\mathbb{I}\{zy < 0\} zy$$

Negative log likelihood (aka cross-entropy)

If z are label probabilities, then

$$\ell(z, y) = -\ln z_y.$$

Regression cost functions

Squared error

If z is a prediction for the dependent variable then

$$\ell(z, y) = (y - z)^2$$

This also corresponds to negative log likelihood under a Gaussianity assumption.

Huber loss

If z is a prediction, then

$$\ell(z, y) = \begin{cases} \frac{1}{2}(z - y)^2 & |z - y| \leq \delta \\ \delta(|z - y| - \frac{1}{2}\delta) & \text{otherwise.} \end{cases} \quad (3)$$

Features and layers

Algorithms

- Random projection

- Back propagation

- Derivatives

- Cost functions

Python libraries

- sklearn

- PyTorch

- TensorFlow

sklearn neural networks

Classification

Uses the **cross entropy** cost

```
from sklearn.neural_network import MLPClassifier
clf = MLPClassifier(hidden_layer_sizes=(5, 2))
clf.fit(X, y)
clf.predict(X_test)
```

- ▶ Main condition is layer sizes.

Regression

```
from sklearn.neural_network import MLPRegressor
model = MLPRegressor(hidden_layer_sizes=(5, 2))
```

Datasets

```
X_train = torch.tensor(X_train, dtype=torch.float32)
train_dataset = TensorDataset(X_train, y_train)
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
mlp = nn.Sequential(nn.Linear(input_size , 50), nn.ReLU())
```

TensorFlow

This is another library, no need to use this for this course