

Bayesian Inference and Hypothesis Testing

Christos Dimitrakakis

October 24, 2024

Conditional Probability and the Theorem of Bayes

Simple Bayesian hypothesis testing

Bayes theorem

- Recall the definition of Conditional probability:

$$P(A|B) = P(A \cap B) / P(B)$$

i.e. the probability of A happening if B happens.

Bayes theorem

- Recall the definition of Conditional probability:

$$P(A|B) = P(A \cap B)/P(B)$$

i.e. the probability of A happening if B happens.

- It is also true that:

$$P(B|A) = P(A \cap B)/P(A)$$

Bayes theorem

- ▶ Recall the definition of Conditional probability:

$$P(A|B) = P(A \cap B)/P(B)$$

i.e. the probability of A happening if B happens.

- ▶ It is also true that:

$$P(B|A) = P(A \cap B)/P(A)$$

- ▶ Combining the two equations, reverse the conditioning:

$$P(A|B) = P(B|A)P(A)/P(B)$$

Bayes theorem

- ▶ Recall the definition of Conditional probability:

$$P(A|B) = P(A \cap B)/P(B)$$

i.e. the probability of A happening if B happens.

- ▶ It is also true that:

$$P(B|A) = P(A \cap B)/P(A)$$

- ▶ Combining the two equations, reverse the conditioning:

$$P(A|B) = P(B|A)P(A)/P(B)$$

- ▶ So we can reverse the order of conditioning, i.e. relate to the probability of A given B to that of B given A.

The cards problem

1. Print out a number of cards, with either $[A|A]$, $[A|B]$ or $[B|B]$ on their sides.
2. If you have an A, what is the probability of an A on the other side?
3. Have the students perform the experiment with:
 - 3.1 Draw a random card.
 - 3.2 Count the number of people with A.
 - 3.3 What is the probability that somebody with an A on one side will have an A on the other?
 - 3.4 Half of the people should have an A?

The cards problem

1. Print out a number of cards, with either $[A|A]$, $[A|B]$ or $[B|B]$ on their sides.
2. If you have an A, what is the probability of an A on the other side?
3. Have the students perform the experiment with:
 - 3.1 Draw a random card.
 - 3.2 Count the number of people with A.
 - 3.3 What is the probability that somebody with an A on one side will have an A on the other?
 - 3.4 Half of the people should have an A?

The prior and posterior probabilities

A	A	2/6	A observed	2/3
A	B	1/6	A observed	1/3
B	A	1/6		
B	B	2/6		

Conditional Probability and the Theorem of Bayes

Simple Bayesian hypothesis testing

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.
- ▶ There are two possibilities:

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.
- ▶ There are two possibilities:
 - ▶ H_0 : They are innocent.

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.
- ▶ There are two possibilities:
 - ▶ H_0 : They are innocent.
 - ▶ H_1 : They are guilty.

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.
- ▶ There are two possibilities:
 - ▶ H_0 : They are innocent.
 - ▶ H_1 : They are guilty.

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.
- ▶ There are two possibilities:
 - ▶ H_0 : They are innocent.
 - ▶ H_1 : They are guilty.

What is your belief that they have committed the crime?

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.
- ▶ There are two possibilities:
 - ▶ H_0 : They are innocent.
 - ▶ H_1 : They are guilty.

What is your belief that they have committed the crime?

Prior elicitation

- ▶ All those that think the accused is guilty, raise their hand.

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.
- ▶ There are two possibilities:
 - ▶ H_0 : They are innocent.
 - ▶ H_1 : They are guilty.

What is your belief that they have committed the crime?

Prior elicitation

- ▶ All those that think the accused is guilty, raise their hand.
- ▶ Divide by the number of people in class

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.
- ▶ There are two possibilities:
 - ▶ H_0 : They are innocent.
 - ▶ H_1 : They are guilty.

What is your belief that they have committed the crime?

Prior elicitation

- ▶ All those that think the accused is guilty, raise their hand.
- ▶ Divide by the number of people in class
- ▶ Let us call this $P(H_1)$.

The murder problem

- ▶ A murder occurred in a house over Christmas. There were n people inside, plus the victim. Person X, the victim's son, is accused of a murder.
- ▶ There are two possibilities:
 - ▶ H_0 : They are innocent.
 - ▶ H_1 : They are guilty.

What is your belief that they have committed the crime?

Prior elicitation

- ▶ All those that think the accused is guilty, raise their hand.
- ▶ Divide by the number of people in class
- ▶ Let us call this $P(H_1)$.
- ▶ This is a purely subjective measure!

DNA test

- ▶ Let us now do a DNA test on the suspect

DNA test

- ▶ Let us now do a DNA test on the suspect

DNA test properties

- ▶ D : Test is positive

DNA test

- ▶ Let us now do a DNA test on the suspect

DNA test properties

- ▶ D : Test is positive
- ▶ $P(D|H_0) = 1\%$: False positive rate

DNA test

- ▶ Let us now do a DNA test on the suspect

DNA test properties

- ▶ D : Test is positive
- ▶ $P(D|H_0) = 1\%$: False positive rate
- ▶ $P(D|H_1) = 100\%$: True positive rate

DNA test

- ▶ Let us now do a DNA test on the suspect

DNA test properties

- ▶ D : Test is positive
- ▶ $P(D|H_0) = 1\%$: False positive rate
- ▶ $P(D|H_1) = 100\%$: True positive rate

DNA test

- ▶ Let us now do a DNA test on the suspect

DNA test properties

- ▶ D : Test is positive
- ▶ $P(D|H_0) = 1\%$: False positive rate
- ▶ $P(D|H_1) = 100\%$: True positive rate

Run the test

- ▶ The result is either positive or negative ($\neg D$).

DNA test

- ▶ Let us now do a DNA test on the suspect

DNA test properties

- ▶ D : Test is positive
- ▶ $P(D|H_0) = 1\%$: False positive rate
- ▶ $P(D|H_1) = 100\%$: True positive rate

Run the test

- ▶ The result is either positive or negative ($\neg D$).
- ▶ What is your belief **now** that the suspect is guilty?

Everybody is a suspect

- ▶ Run a DNA test on everybody in the house.

Everybody is a suspect

- ▶ Run a DNA test on everybody in the house.
- ▶ What is different from before?

Everybody is a suspect

- ▶ Run a DNA test on everybody in the house.
- ▶ What is different from before?
- ▶ Who has a positive test?

Everybody is a suspect

- ▶ Run a DNA test on everybody in the house.
- ▶ What is different from before?
- ▶ Who has a positive test?
- ▶ What is your belief that the people with the positive test are guilty?

Explanation

- ▶ Prior: $P(H_i)$.

Explanation

- ▶ Prior: $P(H_i)$.
- ▶ Likelihood $P(D|H_i)$.

Explanation

- ▶ Prior: $P(H_i)$.
- ▶ Likelihood $P(D|H_i)$.
- ▶ Posterior: $P(H_i|D) = P(D \cap H_i)/P(D) = P(D|H_i)P(H_i)/P(D)$

Explanation

- ▶ Prior: $P(H_i)$.
- ▶ Likelihood $P(D|H_i)$.
- ▶ Posterior: $P(H_i|D) = P(D \cap H_i)/P(D) = P(D|H_i)P(H_i)/P(D)$
- ▶ Marginal probability: $P(D) = P(D|H_0)P(H_0) + P(D|H_1)P(H_1)$

Explanation

- ▶ Prior: $P(H_i)$.
- ▶ Likelihood $P(D|H_i)$.
- ▶ Posterior: $P(H_i|D) = P(D \cap H_i)/P(D) = P(D|H_i)P(H_i)/P(D)$
- ▶ Marginal probability: $P(D) = P(D|H_0)P(H_0) + P(D|H_1)P(H_1)$
- ▶ Posterior: $P(H_0|D) = \frac{P(D|H_0)P(H_0)}{P(D|H_0)P(H_0) + P(D|H_1)P(H_1)}$

Explanation

- ▶ Prior: $P(H_i)$.
- ▶ Likelihood $P(D|H_i)$.
- ▶ Posterior: $P(H_i|D) = P(D \cap H_i)/P(D) = P(D|H_i)P(H_i)/P(D)$
- ▶ Marginal probability: $P(D) = P(D|H_0)P(H_0) + P(D|H_1)P(H_1)$
- ▶ Posterior: $P(H_0|D) = \frac{P(D|H_0)P(H_0)}{P(D|H_0)P(H_0) + P(D|H_1)P(H_1)}$
- ▶ Assuming $P(D|H_1) = 1$, and setting $P(H_0) = q$, this gives

$$P(H_0|D) = \frac{0.1q}{0.1q + 1 - q} = \frac{q}{10 - 9q}$$

Explanation

- ▶ Prior: $P(H_i)$.
- ▶ Likelihood $P(D|H_i)$.
- ▶ Posterior: $P(H_i|D) = P(D \cap H_i)/P(D) = P(D|H_i)P(H_i)/P(D)$
- ▶ Marginal probability: $P(D) = P(D|H_0)P(H_0) + P(D|H_1)P(H_1)$
- ▶ Posterior: $P(H_0|D) = \frac{P(D|H_0)P(H_0)}{P(D|H_0)P(H_0) + P(D|H_1)P(H_1)}$
- ▶ Assuming $P(D|H_1) = 1$, and setting $P(H_0) = q$, this gives

$$P(H_0|D) = \frac{0.1q}{0.1q + 1 - q} = \frac{q}{10 - 9q}$$

- ▶ The posterior can always be updated with more data!

Python example

```
# the input to the function is the prior, the likelihood function, a
# Input:
# - prior for hypothesis 0 (scalar)
# - data (single data point)
# - likelihood[data][hypothesis] array unction
# Returns:
# - posterior for the data point (if multiple points are given, the
def get_posterior(prior, data, likelihood):
    marginal = prior * likelihood[data][0] + (1 - prior) * likelihood[data][1]
    posterior = prior * likelihood[data][0] / marginal
    return posterior

import numpy as np
prior = 0.9
likelihood = np.zeros([2, 2])
# pr of negative test if not a match
likelihood[0][0] = 0.9
# pr of positive test if not a match
likelihood[1][0] = 0.1
# pr of negative test if a match
likelihood[0][1] = 0.1
# pr of positive test if a match
likelihood[1][1] = 0.9
```

Types of hypothesis testing problems

Simple Hypothesis Test

Example: DNA evidence, Covid tests

- ▶ Two hypotheses H_0, H_1
- ▶ $P(D|H_i)$ is defined for all i

Multiple Hypotheses Test

Example: Model selection

- ▶ H_i : One of many mutually exclusive models
- ▶ $P(D|H_i)$ is defined for all i

Null Hypothesis Test

Example: Are men's and women's heights the same?

- ▶ H_0 : The 'null' hypothesis
- ▶ $P(D|H_0)$ is defined
- ▶ The alternative is **undefined**

Pitfalls

Problem definition

- ▶ Defining the models $P(D|H_i)$ incorrectly.

The garden of many paths

Pitfalls

Problem definition

- ▶ Defining the models $P(D|H_i)$ incorrectly.
- ▶ Using an "unreasonable" prior $P(H_i)$

The garden of many paths

Pitfalls

Problem definition

- ▶ Defining the models $P(D|H_i)$ incorrectly.
- ▶ Using an "unreasonable" prior $P(H_i)$

The garden of many paths

- ▶ Having a huge hypothesis space

Pitfalls

Problem definition

- ▶ Defining the models $P(D|H_i)$ incorrectly.
- ▶ Using an "unreasonable" prior $P(H_i)$

The garden of many paths

- ▶ Having a huge hypothesis space
- ▶ Selecting the relevant hypothesis after seeing the data

Bayesian Inference

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$

Bayesian Inference

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Each model β assigns probabilities $P_\beta(x)$ to possible $x \in X$.

Bayesian Inference

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Each model β assigns probabilities $P_\beta(x)$ to possible $x \in X$.
- ▶ We also have a (subjective) prior distribution P over the parameters.

Bayesian Inference

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Each model β assigns probabilities $P_\beta(x)$ to possible $x \in X$.
- ▶ We also have a (subjective) prior distribution P over the parameters.
- ▶ Given x , we calculate the posterior distribution

Bayesian Inference

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Each model β assigns probabilities $P_\beta(x)$ to possible $x \in X$.
- ▶ We also have a (subjective) prior distribution P over the parameters.
- ▶ Given x , we calculate the posterior distribution

Bayesian Inference

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Each model β assigns probabilities $P_\beta(x)$ to possible $x \in X$.
- ▶ We also have a (subjective) prior distribution P over the parameters.
- ▶ Given x , we calculate the posterior distribution

$$P(\beta|x) = \frac{P_\beta(x)P(\beta)}{\sum_{\beta' \in \mathcal{B}} P_{\beta'}(x)P(\beta')}, \quad (\text{finite } \mathcal{B}, P \text{ is a probability})$$

$$p(\beta|x) = \frac{P_\beta(x)p(\beta)}{\int_{\mathcal{B}} P_{\beta'}(x)p(\beta')d\beta'}, \quad (\text{continuous } \mathcal{B}, p \text{ is a density})$$

$$P(B|x) = \frac{\int_B P_{\beta'}(x)dP(\beta)}{\int_{\mathcal{B}} P_{\beta'}(x)dP(\beta)}, \quad B \subset \mathcal{B} \quad (\text{arbitrary } \mathcal{B}, P \text{ is a measure})$$

Bayesian Inference

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Each model β assigns probabilities $P_\beta(x)$ to possible $x \in X$.
- ▶ We also have a (subjective) prior distribution P over the parameters.
- ▶ Given x , we calculate the posterior distribution

$$P(\beta|x) = \frac{P_\beta(x)P(\beta)}{\sum_{\beta' \in \mathcal{B}} P_{\beta'}(x)P(\beta')}, \quad (\text{finite } \mathcal{B}, P \text{ is a probability})$$

$$p(\beta|x) = \frac{P_\beta(x)p(\beta)}{\int_{\mathcal{B}} P_{\beta'}(x)p(\beta')d\beta'}, \quad (\text{continuous } \mathcal{B}, p \text{ is a density})$$

$$P(B|x) = \frac{\int_B P_{\beta'}(x)dP(\beta)}{\int_{\mathcal{B}} P_{\beta'}(x)dP(\beta)}, \quad B \subset \mathcal{B} \quad (\text{arbitrary } \mathcal{B}, P \text{ is a measure})$$

Alternative notation for different probability spaces

- ▶ The **prior** $P(\beta) = \mathbb{P}(\beta)$ and **posterior** $P(\beta | x) = \mathbb{P}(\beta | x)$ belief.
- ▶ The **likelihood** $P_\beta(x) = \mathbb{P}(x | \beta)$
- ▶ The **marginal** $\mathbb{P}_P(x) = \sum_{\beta} P_\beta(x)P(\beta)$.

Probabilistic machine learning

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$

Probabilistic machine learning

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Prior P on \mathcal{B}

Probabilistic machine learning

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Prior P on \mathcal{B}
- ▶ Observations $x = x_1, \dots, x_t$.

Probabilistic machine learning

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Prior P on \mathcal{B}
- ▶ Observations $x = x_1, \dots, x_t$.

Probabilistic machine learning

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Prior P on \mathcal{B}
- ▶ Observations $x = x_1, \dots, x_t$.

Maximum likelihood approach

- ▶ Model selection: $\beta_{ML}^*(x) = \arg \max_{\beta} P_\beta(x)$.
- ▶ Model prediction: $P_{\beta_{ML}^*(x)}(x_{t+1})$

Probabilistic machine learning

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Prior P on \mathcal{B}
- ▶ Observations $x = x_1, \dots, x_t$.

Maximum likelihood approach

- ▶ Model selection: $\beta_{ML}^*(x) = \arg \max_{\beta} P_\beta(x)$.
- ▶ Model prediction: $P_{\beta_{ML}^*(x)}(x_{t+1})$

Maximum a posteriori approach

- ▶ Model selection: $\beta_{MAP}^*(x) = \arg \max_{\beta} P_\beta(x)P(\beta)$.
- ▶ Model prediction: $P_{\beta_{MAP}^*(x)}(x_{t+1})$

Probabilistic machine learning

- ▶ Model family $\{P_\beta | \beta \in \mathcal{B}\}$
- ▶ Prior P on \mathcal{B}
- ▶ Observations $x = x_1, \dots, x_t$.

Maximum likelihood approach

- ▶ Model selection: $\beta_{ML}^*(x) = \arg \max_{\beta} P_\beta(x)$.
- ▶ Model prediction: $P_{\beta_{ML}^*(x)}(x_{t+1})$

Maximum a posteriori approach

- ▶ Model selection: $\beta_{MAP}^*(x) = \arg \max_{\beta} P_\beta(x)P(\beta)$.
- ▶ Model prediction: $P_{\beta_{MAP}^*(x)}(x_{t+1})$

Bayesian approach

- ▶ Posterior calculation: $P(\beta|x) = P_\beta(x)P(\beta)/\mathbb{P}_P(x)$
- ▶ Model prediction: $\mathbb{P}_P(x_{t+1}|x) = \sum_{\beta} P_\beta(x_{t+1})P(\beta|x)$

Differences between approaches

Maximum likelihood approach

- ▶ Ignores model complexity
- ▶ Is an optimisation problem

Maximum a posteriori approach

- ▶ Regularises model selection using the prior
- ▶ Can be seen as solving the optimisation problem

$$\max_{\beta} \ln P_{\beta}(x) + \ln P(\beta),$$

where the prior term $\ln P(\beta)$ acts as a regulariser.

Bayesian approach

- ▶ Does not select a single model
- ▶ Averages over all models according to their fit **and** the prior
- ▶ Does **not** result in an optimisation problem.

The n-meteorologists problem

- ▶ Consider n meteorological stations $\{\mu\}$ predicting rainfall.
- ▶ $x_t \in \{0, 1\}$ with $x_t = 1$ if it rains on day t .
- ▶ We have a prior distribution $P(\mu)$ for each station.
- ▶ At time t , station μ makes as a prediction $P_\mu(x_{t+1}|x_1, \dots, x_t)$
- ▶ We observe x_{t+1} and calculate the posterior $P(\mu|x_1, \dots, x_t, x_{t+1})$.

The marginal distribution

To take into account all stations, we can marginalise:

$$\mathbb{P}_P(x_{t+1} \mid x_1, \dots, x_t) = \sum_{\mu} P_\mu(x_{t+1}|x_t)P(\mu)$$

The posterior

- ▶ Show that

$$P(\mu \mid x_1, \dots, x_{t+1}) = \frac{P_\mu(x_t \mid x_1, \dots, x_t)P(\mu|x_1, \dots, x_t)}{\sum_{\mu'} P_{\mu'}(x_t \mid x_1, \dots, x_t)P(\mu'|x_1, \dots, x_t)}$$

- ▶ How would you implement an ML or a MAP solution to this problem?

Sufficient statistics

A statistic f

This is any function $f : X \rightarrow S$ where

- ▶ X is the data space
- ▶ S is an arbitrary space

Example statistics for $X = \mathbb{R}^*$ (the set of all real-valued sequences)

- ▶ The sample mean of a sequence $1/T \sum_{t=1}^T x_t$
- ▶ The total number of samples T

Sufficient statistic

f is sufficient for a family $\{P_\beta : \beta \in \mathcal{B}\}$ when

$$f(x) = f(x') \Rightarrow P_\beta(x) = P_\beta(x') \forall \beta \in \mathcal{B}.$$

If there exists a finite-dimensional sufficient statistic, Bayesian and ML learning can be done in closed form within the family.

Conjugate priors

Consider a parametrised family of priors \mathcal{P} on \mathcal{B} and a distribution family $\{P_\beta\}$. The pair is conjugate if, for any prior $P \in \mathcal{P}$, and any observation x , there exists $P' \in \mathcal{P}$ such that $P'(\beta) = P(\beta|x)$.

Standard Parametric conjugate families

Prior	Likelihood	Parameters β	Observations x
Beta	Bernoulli	$[0, 1]$	$\{0, 1\}^T$
Multinomial	Dirichlet	n	$\{1, \dots, n\}^T$
Gamma	Normal	\mathbb{R}, \mathbb{R}	\mathbb{R}^T
Wishart	Normal	$\mathbb{R}^n, \mathbb{R}^{n \times n}$	$\mathbb{R}^{n \times T}$

The Simplex $n = \{\beta \in [0, 1]^n : \|\beta\|_1 = 1\}$ is the set of all n -dimensional probability vectors.

Extensions

- ▶ Discrete Bayesian Networks.
- ▶ Linear-Gaussian Models (i.e. Bayesian linear regression)
- ▶ Gaussian Processes.

Beta-Bernoulli



Definition of the Bernoulli distribution

If $x_t \mid \beta \sim \text{Bernoulli}(\beta)$. $\beta \in [0, 1]$, $x_t \in \{0, 1\}$ and:

$$P_\beta(x_t = 1) = \beta$$

Definition of the Beta density

If $\beta \sim \text{Beta}(\alpha_1, \alpha_0)$, $\alpha_0, \alpha_1 > 0$ and

$$p(\beta \mid \alpha_1, \alpha_0) \propto \beta^{\alpha_1 - 1} (1 - \beta)^{\alpha_0 - 1}$$

Beta-Bernoulli conjugate pair

- ▶ $\beta \sim \text{Beta}(\alpha_1, \alpha_0)$.
- ▶ $x_t \mid \beta \sim \text{Bernoulli}(\beta)$.

Then, for any $x = x_1, \dots, x_T$, the posterior distribution is

- ▶ $\beta \mid x \sim \text{Beta}(\alpha_1 + \sum_t x_t, \alpha_0 + T - \sum_t x_t)$.

Dirichlet-Multinomial



Definition of the Multinomial distribution

If $x_t \mid \beta \sim (\beta)$, with $\beta \in^n$ and $x_t \in \{1, \dots, n\}$ and:

$$P_{\beta}(x_t = i) = \beta_i$$

Definition of the Dirichlet density

If $\beta \sim (\alpha)$, with $\alpha \in \mathbb{R}_+^n$ then

$$p(\beta|\alpha) \propto \prod_i \beta_i^{\alpha_i-1}$$


Dirichlet-Multinomial conjugate pair

- ▶ $\beta \sim (\alpha)$.
- ▶ $x_t \mid \beta \sim \text{Bernoulli}(\beta)$.

Then, for any $x = x_1, \dots, x_T$, the posterior distribution is


- ▶ $\beta \mid x \sim (\alpha + s_T)$, where $s_T = \sum_{t=1}^T \mathbf{1}_{x_t = i}$

Discrete Bayesian Networks



- ▶ A directed acyclic graph (DAG) defined on variables x_1, \dots, x_n with each x_n taking a finite number of values,
- ▶ Let S_i be the indices corresponding to parent variables of x_i .
- ▶ $x_i \mid \beta_i, x_{S_i} = k \sim (\beta_{i,k})$.

Example: Lung cancer, smoking and asbestos




$$P_{\beta_A}(x_A = 1) = \beta_A \quad (1)$$

$$P_{\beta_S}(x_S = 1) = \beta_S \quad (2)$$

$$P_{\beta_C}(x_C = 1 \mid X_A = j, X_S = k) = \beta_{C,j,k} \quad (3)$$

Markov model



A **Markov model** obeys

$$\mathbb{P}_{\beta}(x_{k+1} | x_k, \dots, x_1) = \mathbb{P}_{\beta}(x_{k+1} | x_k)$$

i.e. the graphical model is a chain. We are usually interested in **homogeneous** models, where

$$\mathbb{P}_{\beta}(x_{k+1} = i | x_k = j) = \beta_{i,j} \quad \forall k$$

Inference for finite Markov models

- ▶ If $x_t \in [n]$ then $x_{t+1} | \beta, x_t = i \sim (\beta_i)$, $\beta_i \in^n$
- ▶ Prior $\beta_i | \alpha \sim (\alpha)$ for all $i \in [n]$.
- ▶ Posterior $\beta_i | x_1, \dots, x_t, \alpha \sim (\alpha^{(t)})$ with

$$\alpha_{i,j}^t = \alpha_{i,j} + \sum_{k=1}^t \mathbb{I}\{x_k = i \wedge x_{k+1} = j\}, \quad \alpha^0 = \alpha.$$