

# Generative Modelling

Christos Dimitrakakis

October 31, 2023

# Outline

## Classification

- Classification: Generative modelling
- Density estimation

## Algorithms for latent variable models

- Gradient algorithms
- Expectation maximisation

## Exercises

- Density estimation
- Classification

## Classification

Classification: Generative modelling

Density estimation

Algorithms for latent variable models

Exercises

# Generative modelling

## general idea

- ▶ Data  $(x, y)$ .
- ▶ Need to model  $P(y|x)$ .
- ▶ Model the complete data distribution:  $P(x|y)$ ,  $P(x)$ ,  $P(y)$ .
- ▶ Calculate  $P(y|x) = \frac{P(x|y)P(y)}{P(x)}$ .

## Examples


- ▶ Naive Bayes classifier
- ▶ Gaussian Mixture Classifier

## Modelling the data distribution


- ▶ Need to estimate the density  $P(x|y)$  for each class  $y$ .
- ▶ Need to estimate  $P(y)$

# The basic graphical model


## A discriminative classification model

Here  $P(y|x)$  is given directly. 

## A generative classification model


Here  $P(y|x) = P(x|y)P(y)/P(x)$ . 

## A generative model

Here we just have  $P(x)$ . 


# Adding parameters to the graphical model

## A Bernoulli RV



Here,  $x|\theta \sim \text{Bernoulli}(\theta)$

## A normally distributed variable



Here  $x|\mu, \sigma \sim \text{Normal}(\mu, \sigma^2)$

# Classification: Naive Bayes Classifier

- ▶ Data  $(x, y)$
- ▶  $x \in X$
- ▶  $y \in Y \subset \mathbb{N}$ ,  $N_i$ : amount of data from class  $i$ .

## Separately model each class

- ▶ Assume each class data comes from a different normal distribution
- ▶  $x|y = i \sim \text{Normal}(\mu_i, \sigma_i I)$
- ▶ For each class, calculate
  - ▶ Empirical mean  $\hat{\mu}_i = \sum_{t: y_t = i} x_t / N_i$
  - ▶ Empirical variance  $\hat{\sigma}_i$ .

## Decision rule

Use Bayes's theorem:

$$P(y|x) = P(x|y)P(y)/P(x),$$

choosing the  $y$  with largest posterior  $P(y|x)$ .

- ▶  $P(x|y = i) \propto \exp(-\|\hat{\mu}_i - x\|^2 / \hat{\sigma}_i^2)$

# Graphical model for the Naive Bayes Classifier


When  $x \in \mathbb{R}$

Assume  $k$  classes, then

►  $\mu = (\mu_1, \dots, \mu_k)$

►  $\sigma = (\sigma_1, \dots, \sigma_k)$

►  $\theta = (\theta_1, \dots, \theta_k)$



►  $y \mid \theta \sim \text{Multinomial}(\theta)$

►  $x \mid y, \mu, \sigma \sim \text{Normal}(\mu_y, \sigma_y^2)$



# General idea

## Parametric models

- ▶ Fixed histograms
- ▶ Gaussian Mixtures

## Non-parametric models

- ▶ Variable-bin histograms
- ▶ Infinite Gaussian Mixture Model
- ▶ Kernel methods

# Histograms

## Fixed histogram

- ▶ Hyper-Parameters: number of bins
- ▶ Parameters: Number of points in each bin.

## Variable histogram

- ▶ Hyper-parameters: Rule for constructing bins
- ▶ Generally  $\sqrt{n}$  points in each bin.

# Gaussian Mixture Model

## Hyperparameters:

- ▶ Number of Gaussian  $k$ .

## Parameters:

- ▶ Multinomial distribution  $\theta$  over Gaussians
- ▶ For each Gaussian  $i$ , center  $\mu_i$ , covariance matrix  $\Sigma_i$ .

## Model. For each point $x_t$ :

- ▶  $c_t = i$  w.p.  $\theta_i$
- ▶  $x_t | c_t = i \sim \text{Normal}(\mu_i, \Sigma_i)$ .

## Algorithms:

- ▶ Expectation Maximisation
- ▶ Gradient Ascent
- ▶ Variational Bayesian Inference (with appropriate prior)

# Gradient ascent

## Objective function

$$L(\theta) = P(x|\theta)$$

## Marginalisation over latent variable

$$L(\theta) = \sum_z P(z, x | \theta)$$

## Gradient ascent

$$\theta^{(n+1)} = \theta^{(n)} + \alpha \nabla_{\theta} L(\theta)$$

## Gradient calculation

Here we use the **log trick**:  $\nabla \ln f(x) = \nabla f(x)/f(x)$ .

$$\nabla_{\theta} L(\theta) = \sum_z \nabla_{\theta} P(z, x | \theta) \quad (1)$$

$$= \sum_z P(z, x | \theta) \nabla_{\theta} \ln P(z, x | \theta) \quad (2)$$

$$= \sum_z P(x | z, \theta) P(z | \theta) \nabla_{\theta} \ln P(z, x | \theta) \quad (3)$$

$$\approx \frac{1}{m} \sum_{i=1}^m P(x | z^{(i)}, \theta) \nabla_{\theta} \ln P(z^{(i)}, x | \theta) \quad z^{(i)} \sim P(z | \theta) \quad (4)$$

## A lower bound on the likelihood

$$\begin{aligned}\ln P(x|\theta) &= \sum_z G(z) P(x|\theta) \\&= \sum_z G(z) [\ln P(x, z|\theta) - \ln P(z|x, \theta)] \\&= \sum_z G(z) \ln P(x, z|\theta) - \sum_z G(z) \ln P(z|x, \theta) \\&= \sum_z P(z|x, \theta^{(k)}) \ln P(x, z|\theta) - \sum_z P(z|x, \theta^{(k)}) \ln P(z|x, \theta) \\&\geq \sum_z P(z|x, \theta^{(k)}) \ln P(x, z|\theta) - \sum_z P(z|x, \theta^{(k)}) \ln P(z|x, \theta^{(k)}) \\&= Q(\theta | \theta^{(k)}) + \mathbb{H}(z | x = x, \theta = \theta^{(k)})\end{aligned}$$

## The Gibbs Inequality

$D_{KL}(P\|Q) \geq 0$ , or  $\sum_x \ln P(x)P(x) \geq \sum_x \ln Q(x)P(x)$ .

# EM Algorithm (Dempster et al, 1977)

- ▶ Initial parameter  $\theta^{(0)}$ , observed data  $x$
- ▶ For  $k = 0, 1, \dots$

– Expectation step:

$$Q(\theta \mid \theta^{(k)}) \triangleq \mathbb{E}_{z \sim P(z|x, \theta^{(k)})} [\ln P(x, z | \theta)] = \sum_z [\ln P(x, z | \theta)] P(z \mid x, \theta^{(k)})$$

– Maximisation step:

$$\theta^{(k+1)} = \arg \max_{\theta} Q(\theta, \theta^{(k)}).$$

See *Expectation-Maximization as lower bound maximization*, Minka, 1998

# Minorise-Maximise

EM can be seen as a version of the minorise-maximise algorithm

- ▶  $f(\theta)$ : Target function to **maximise**
- ▶  $Q(\theta|\theta^{(k)})$ : surrogate function

## $Q$ Minorizes $f$

This means surrogate is always a lower bound so that

$$f(\theta) \geq Q(\theta|\theta^{(k)}), \quad f(\theta^{(k)}) \geq Q(\theta^{(k)}|\theta^{(k)}),$$

## Algorithm

- ▶ Calculate:  $Q(\theta|\theta^{(k)})$
- ▶ Optimise:  $\theta^{(k+1)} = \arg \max_{\theta} Q(\theta|\theta^{(k)})$ .

# GMM versus histogram

- ▶ Generate some data  $x$  from an arbitrary distribution in  $\mathbb{R}$ .
- ▶ Fit the data with a histogram for varying numbers of bins
- ▶ Fit a GMM with varying numbers of Gaussians
- ▶ What is the best fit? How can you measure it?



# GMM Classifier

## Base class: sklearn GaussianMixtureModel

- ▶ *fit()* only works for Density Estimation
- ▶ *predict()* only predicts cluster labels

## Problem

- ▶ Create a GMMClassifier class
- ▶ *fit()* should take X, y, arguments
- ▶ *predict()* should predict class labels
- ▶ Hint: Use *predict\_proba()* and multiple GMM models