

# Nearest Neighbour Algorithms

Christos Dimitrakakis

September 20, 2024

# Outline

## Introduction

The hidden secret of machine learning

## The algorithm

$k$  Nearest Neighbours

Extensions and parameters

## Activities

## Introduction

The hidden secret of machine learning

## The algorithm

$k$  Nearest Neighbours

Extensions and parameters

## Activities

# Supervised learning

- ▶ Given labelled training examples  $(x_1, y_1), \dots, (x_T, y_T)$  where
- ▶  $x_t \in X$ , the **input** variables (or **features**)
- ▶  $y_t \in Y$ , the **output** variables (or **labels**)

## Feature space $\mathcal{X}$

- ▶ Usually  $\mathcal{X} = \mathbb{R}^n$ : the n-dimensional Euclidean space
- ▶ How do we use your class data?

## Classification

- ▶  $Y = \{1, \dots, m\}$  are **discrete** labels

## Regression

- ▶  $Y = \mathbb{R}^m$  are **continuous** values

# The kNN algorithm idea

- ▶ Assume an unknown example is similar to its neighbours
- ▶ Smoothness allows us to make predictions

*Discriminatory analysis-nonparametric discrimination: consistency properties*, Evelyn Fix and Joseph L. Hodges Jr, 1951.

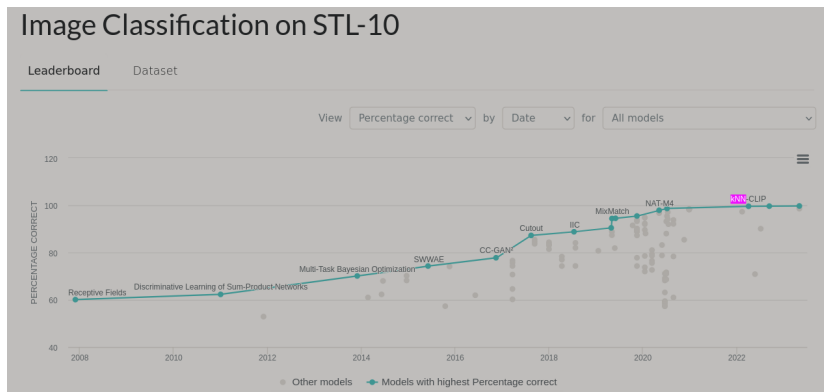


Figure: Evelyn Fix



Figure: Joseph Hodges

# Performance of KNN on image classification



- ▶ Really simple!
- ▶ Can outperform really complex models!

## Introduction

The hidden secret of machine learning

## The algorithm

$k$  Nearest Neighbours

Extensions and parameters

## Activities

# The Nearest Neighbour algorithm

## Pseudocode

- ▶ Input: Data  $(x_t, y_t)_{t=1}^T$ , test point  $x$ , distance  $d$
- ▶  $t^* = \arg \min_t d(x_t, x)$  / How do we implement this?
- ▶ Return  $\hat{y}_t = y_{t^*}$

## Classification

$$\hat{y}_t \in [m] \equiv \{1, \dots, m\}$$

## Regression

$$\hat{y}_t \in \mathbb{R}^m$$



# The $k$ -Nearest Neighbour algorithm

## Pseudocode

- ▶ Input: Data  $(x_t, y_t)_{t=1}^T$ , test point  $x$ , distance  $d$ , neighbours  $k$
- ▶ Calculate  $h_t = d(x_t, x)$  for all  $t$ .
- ▶ Get sorted indices  $s = \text{argsort}(h)$  so that  $d(x_{s_i}, x) \leq d(x_{s_{i+1}}, x)$  for all  $i$ . (How?)
- ▶ Return  $\sum_{i=1}^k y_{s_i} / k$ .

## Classification

- ▶ It is not convenient to work with discrete labels.
- ▶ We use a **one-hot encoding**  $(0, \dots, 0, 1, 0, \dots, 0)$ .
- ▶  $y_t \in \{0, 1\}^m$  with  $\|y_t\|_1 = 1$ , so that the class of the  $t$ -th example is  $j$  iff  $y_{t,j} = 1$ .

## Regression

- ▶  $y_t \in \mathbb{R}^m$ , so we need do nothing

# Making a decision

kNN: A **model** of the conditional distribution  $P(y|x)$

- ▶ Given features  $x$ , we get a vector
- ▶  $p_i = \hat{\mathbb{P}}(y = i|x)$ .

The optimal decision rule  $\pi$  derived from kNN

- ▶ Classification decision  $a_t \sim \pi(a|x_t)$
- ▶  $a_t \in$  but  $\neq \mathcal{Y}$ , e.g. can include "Do not Know", or "Alert" etc.
- ▶ Actual label  $y_t$
- ▶  $U(a_t, y_t)$ : utility function depending on the application.

Decision rule maximising accuracy

- ▶  $a_t = \arg \max_i \hat{\mathbb{P}}(y = i|x)$ .

# The number of neighbours

$$k = 1$$

- ▶ How does it perform on the training data?
- ▶ How might it perform on unseen data?

$$k = T$$

- ▶ How does it perform on the training data?
- ▶ How might it perform on unseen data?

## Distance function

For data in  $\mathbb{R}^n$ ,  $p$ -norm

$$d(x, y) = \|x - y\|_p$$

## Scaled norms

When features having varying scales:

$$d(x, y) = \|Sx - Sy\|_p$$

Or pre-scale the data

## Complex data

- ▶ Manifold distances
- ▶ Graph distance

# Distances

A distance  $d(\cdot, \cdot)$ :

- ▶ Identity  $d(x, x) = 0$ .
- ▶ Positivity  $d(x, y) > 0$  if  $x \neq y$ .
- ▶ Symmetry  $d(y, x) = d(x, y)$ .
- ▶ Triangle inequality  $d(x, y) \leq d(x, z) + d(z, y)$ .

For data in  $\mathbb{R}^n$ ,  $p$ -norm

$$d(x, y) = \|x - y\|_p$$

# Norms;

## A norm $\|\cdot\|$

- ▶ Zero element  $\|0\| = 0$ .
- ▶ Homogeneity  $\|cx\| = c\|x\|$  for any scalar  $a$ .
- ▶ Triangle inequality  $\|x + y\| \leq \|x\| + \|y\|$ .

## $\ell_p$ -norm

$$\|z\|_p = \left( \sum_i z_i^p \right)^{1/p}$$

# Neighbourhood calculation

If we have  $T$  datapoints

Sort and top  $K$ .

- Requires  $O(T \ln T)$  time

Use the Cover-Tree or KD-Tree algorithm

- Requires  $O(cK \ln T)$  time.
- $c$  depends on the data distribution.

## Introduction

The hidden secret of machine learning

## The algorithm

$k$  Nearest Neighbours

Extensions and parameters

## Activities



# KNN activity

- ▶ Implement nearest neighbours
- ▶ Introduction to scikitlearn nearest neighbours
- ▶ Introduction to generalisation errors