

Machine Learning and Data Mining

Christos Dimitrakakis

August 31, 2023

Outline

General

Introduction

Material

The problems of Machine Learning (1 week)

Introduction

Activities

Models, hypotheses

Examples

Pitfalls

Learning as Optimisation (4 weeks)

Objective functions

k Nearest Neighbours

Supervised machine learning problems

Learning and generalisation

Cross-validation lab

activity

Linear neural networks

Multi-layer neural networks

Support Vector Machines

Learning as Probabilistic Inference (4 weeks)

Probabilistic Models

Classification: Discriminative modelling

Summary

This course gives an introduction to the algorithms and theory of machine learning. Application is in the form of a course project. During the course, you will be able to:

- ▶ Formulate machine learning problems in terms of optimisation or probabilistic inference.
- ▶ Understand the fundamental machine learning algorithms.
- ▶ Be able to implement some of the simplest algorithms.
- ▶ Apply off-the-shelf algorithms from scikit-learn to problems.
- ▶ Develop custom models using the pyTorch library.

The course focuses on algorithms and models, firstly on optimisation-based learning, and the secondly on probabilistic learning.

Reference Material

Topic	ISLP	ESL2
Linear Regression	3	3
Nearest Neighbours	3,4	13
Linear Classification	4	4
Model Selection	5	7
Linear Model Regularization	6	3
Basis Expansions	7	5
Kernel Smoothing	7	6
Additive Models	7	9
Model Inference/Averaging	8	8
Random Forests	8	15
Ensemble Learning	8	16
Trees	8	9
Boosting	8	10
SVMs	9	12
Neural Netowrks	10	11
Censored Data	11	18
Unsupervised Learning	12	14
Undirected Graphical Models	*	17
Hypothesis tesing	13	18
High Dimensional Statistics	6	18

Schedule

Week	Date	Topic	Theory
1	09.26	Supervised Learning, kNN	Decision theory
2	10.03	XV, Bootstrapping	Generalisation
3	10.10	Perceptron	SGD, Convergence
4	10.17	Discriminative models	Linear Regression, Least-Sq
5	10.24	Generative Models	Bayes Classifier
6	10.31	Basis Functions	GAMs
7	11.07	Multi-Layer Neural Network	Backpropagation
8	11.14	Support Vector Machines	Maximal Margin
9	11.21	Regularisation	Non-linear programming
10	11.28	Bayesian Inference	Conjugate priors
11	12.05	Latent Variable Models	Expectation Maximisation
12	12.12	Approximate Bayesian Inference	Monte-Carlo Methods
13	12.19	Project Presentations	

Textbooks

Primary

- ▶ Introduction to Statistical Learning with Python

https://hastie.su.domains/ISLP/ISLP_website.pdf

- ▶ Elements of Statistical Learning

<https://hastie.su.domains/Papers/ESLII.pdf>

Secondary

- ▶ Probabilistic Machine Learning: An Introduction

<https://probml.github.io/pml-book/book1.html>

<https://github.com/probml/pml-book/releases/latest/download/book1.pdf>

- ▶ Probabilistic Machine Learning: Advanced Topics

<https://probml.github.io/pml-book/book2.html>

<https://github.com/probml/pml2-book/releases/latest/download/book2.pdf>

General

The problems of Machine Learning (1 week)

- Introduction

- Activities

- Models, hypotheses

- Examples

- Pitfalls

Learning as Optimisation (4 weeks)

Learning as Probabilistic Inference (4 weeks)

Sequence modelling (2 weeks)

Reinforcement Learning (2 weeks)

Machine learning

Data Collection

- ▶ Downloading a clean dataset from a repository
- ▶ Performing a survey
- ▶ Scraping data from the web
- ▶ Deploying sensors, performing experiments, and obtaining measurements.

Modelling (what we focus on this course)

- ▶ Can be as simple as counting coin tosses.
- ▶ Can be as complex as a large language model with billions of parameters
- ▶ The model depends on the data and the problem

Decision Making

- ▶ Ultimately, we use models to make decisions.
- ▶ However, decisions are made every step of the way (how to collect data, which model to choose)

Class data

Fill in your data (does not have to be true)

The main problems in machine learning and statistics

Prediction

- ▶ Will it rain tomorrow?
- ▶ How much will bitcoin be worth next year?

Inference

- ▶ Does my poker opponent have two aces?
- ▶ What is the mass of the moon?
- ▶ What is the law of gravitation?

Decision Making

- ▶ Should I go hiking tomorrow?
- ▶ Should I buy some bitcoins?
- ▶ Should I fold, call, or raise in my poker game?
- ▶ How can I get a spaceship to orbit the moon?

The need to learn from data

Problem definition

- ▶ What problem do we need to solve?
- ▶ How can we formalise it?
- ▶ What properties of the problem can we learn from data?

Data collection

- ▶ Why do we need data?
- ▶ What data do we need?
- ▶ How much data do we want?
- ▶ How will we collect the data?

Modelling and decision making

- ▶ How will we compute something useful?

Learning from data

Unsupervised learning

- ▶ Given data x_1, \dots, x_T .
- ▶ Learn about the data-generating process.

Supervised learning

- ▶ Given data $(x_1, y_1), \dots, (x_T, y_T)$
- ▶ Learn about the relationship between x_t and y_t .
- ▶ Example: Classification, Regression

Online learning

- ▶ Sequence prediction: At each step t , predict x_{t+1} from x_1, \dots, x_t .
- ▶ Conditional prediction: At each step t , predict y_{t+1} from $x_1, y_1, \dots, x_t, y_t, x_{t+1}$

Reinforcement learning

Learn to act in an **unknown** world through interaction and rewards

Unsupervised learning

Image compression

- ▶ Learn two mappings c, d
- ▶ $c(x)$ compresses an image x to a small representation z .
- ▶ $d(z)$ decompresses to an approximate image \hat{x} .

Supervised learning

Image classification

Unsupervised learning

Density estimation

Compression

Generative modelling

Pitfalls

Reproducibility

- ▶ Modelling assumptions
- ▶ Distribution shift
- ▶ Interactions and feedback

Fairness

- ▶ Implicit biases in training data
- ▶ Fair decision rules and meritocracy

Privacy

- ▶ Accidental data disclosure
- ▶ Re-identification risk

General

The problems of Machine Learning (1 week)

Learning as Optimisation (4 weeks)

Objective functions

k Nearest Neighbours

Supervised machine learning problems

Learning and generalisation

Cross-validation lab

Linear neural networks

Multi-layer neural networks

Support Vector Machines

activity

Learning as Probabilistic Inference (4 weeks)

Sequence modelling (2 weeks)

Reinforcement Learning (2 weeks)

Supervised learning objectives

- ▶ Data (x_t, y_t) , $x_t \in X$, $y_t \in Y$, $t \in [T]$.
- ▶ i.i.d assumption: $(x_t, y_t) \sim P$ for all t .
- ▶ Supervised decision rule $\pi(a_t|x_t)$

Classification

- ▶ Predict the labels correctly, i.e. $a_t = y_t$.
- ▶ Have an appropriate confidence level

Regression

- ▶ Predict the mean correctly
- ▶ Have an appropriate variance around the mean

Unsupervised learning objectives

- ▶ Reconstruct the data well
- ▶ Model the data-generating distribution
- ▶ Be able to generate data

Reinforcement learning objectives

- ▶ Maximise total expected reward, either
- ▶ during learning, or
- ▶ after learning is finished.

A simple classification problem

Height distribution data:

- ▶ $y \in \{M, F\}$, gender.
- ▶ $x \in \mathbb{R}$, income.

Problems

- ▶ Can we model the height distribution $P(x)$?
- ▶ $P(x|y)$ How does the height depend on the gender?
- ▶ $P(y|x)$ How does the gender depend on the height?

The Bayes classifier

- ▶ Predicted gender a from height so that
 $a = \arg \max_y P(y|x)$.
- ▶ Requires knowledge of P .

The Nearest Neighbour algorithm

Pseudocode

- ▶ Input: Data $(x_t, y_t)_{t=1}^T$, test point x , distance d
- ▶ $t^* = \arg \min_t d(x_t, x)$
- ▶ Return $y^* = y_{t^*}$

Classification

$y_t \in [m] \equiv \{1, \dots, m\}$ See example code

Regression

$y_t \in \mathbb{R}^m$

The k-Nearest Neighbour algorithm

Pseudocode

- ▶ Input: Data $(x_t, y_t)_{t=1}^T$, test point x , distance d , neighbours k
- ▶ Calculate $h_t = d(x_t, x)$ for all t .
- ▶ Get sorted indices $s = \text{argsort}(h)$ so that $d(x_{s_i}, x) \leq d(x_{s_{i+1}}, x)$ for all i .
- ▶ Return $\sum_{i=1}^k y_{s_i} / k$.

Classification

- ▶ It is not convenient to work with discrete labels
- ▶ We use a **one-hot encoding** vector representation $(0, \dots, 0, 1, 0, \dots, 0)$.
- ▶ $y_t \in \{0, 1\}^m$ with $\|y_t\|_1 = 1$, so that the class of the t -th example is j iff $y_{t,j} = 1$.

Regression

$$y_t \in \mathbb{R}^m$$

Code:

Classification

The classifier as a decision rule

A decision rule $\pi(a|x)$ generates a **decision** $a \in [m]$. It is the conditional probability of a given x .

Even though normally conditional probabilities are defined as $P(A|B) = P(A \cap B)/P(B)$, the probability of the decision a is undefined without a given x . So it's better to

The accuracy of a single decision

$$U(a_t, y_t) = \mathbb{I}\{a_t = y_t\} = \begin{cases} 1, & \text{if } a_t = y_t \\ 0, & \text{otherwise} \end{cases}$$

$$U(\pi, D) \triangleq \frac{1}{T} \sum_{t=1}^T \sum_{a=1}^m \pi(y_t|x_t)$$

The accuracy on the training set

$$U(\pi, D) \triangleq \frac{1}{T} \sum_{t=1}^T \sum_{a=1}^m \pi(y_t|x_t)$$

Regression

The regressor as a decision rule

A decision rule $\pi(a|x)$ generates a **decision** $a \in \mathbb{R}^m$. It is the conditional density of a given x .

Accuracy

If $(x, y) \sim P$, the accuracy U of a decision rule π under the distribution P is:

$$U(\pi, P) \triangleq \int_X \int_Y dP(x, y) \pi(y|x).$$

Mean-Squared Error

If $(x, y) \sim P$, the mean-square error of a deterministic decision rule $\pi : X \rightarrow \mathbb{R}$ under the distribution $P(x, y) = P(x|y)P(y)$ is:

$$\int_X \sum_{y=1}^m dP(x|y)P(y) \sum_{a=1}^m \pi(a|x)$$

The Train/Test methodology

Training data $D = ((x_t, y_t) : t = 1, \dots, T)$.

- ▶ $x_t \in X$
- ▶ $y_t \in \mathbb{R}^m$.

Assumption: The data is generated i.i.d.

- ▶ $(x_t, y_t) \sim P$ for all t (identical)
- ▶ $D \sim P^T$ (independent)

The optimal decision rule for P

$$\max_{\pi} U(\pi, P) = \max_{\pi} \int_{X \times Y} dP(x, y) \sum_a \pi(a|x) U(a, y)$$

The optimal decision rule for D

$$\max_{\pi} U(\pi, D) = \max_{\pi} \sum_{(x,y) \in D} \sum_a \pi(a|x) U(a, y)$$

Generalisation as error

Error due to mismatched objectives

The π^* maximising $U(\pi, P)$ is not the $\hat{\pi}$ maximising $U(\pi, D)$.

Lemma

If $|U(\pi, P) - U(\pi, D)| \leq \epsilon$ for all π then

$$U(\hat{\pi}, D) \geq U(\pi^*, P) - 2\epsilon.$$

Error due to restricted classes

- ▶ We may use a constrained $\hat{\Pi} \subset \Pi$.
- ▶ Then $\max_{\hat{\pi} \in \hat{\Pi}} U(\pi, P) \leq \max_{\pi \in \Pi} U(\pi, P)$.

The bias/variance trade-off

- ▶ Dataset D P .
- ▶ Predictor $f_D(x)$
- ▶ Target function $y = f(x) + \epsilon$
- ▶ $\mathbb{E} \epsilon = 0$ zero-mean noise with variance $\sigma^2 = \mathbb{V}(\epsilon)$

MSE decomposition

$$\mathbb{E}[(f - f_D)^2] = \mathbb{V}(f_D) + \mathbb{B}(f_D)^2 + \sigma^2$$

Variance

How sensitive the estimator is to the data

$$\mathbb{V}(f_D) = \mathbb{E}[(f_D - \mathbb{E}(f_D))^2]$$

Bias

What is the expected deviation from the true function

$$\mathbb{B}(f_D) = \mathbb{E}[(f_D - f)]$$

Example: mean estimation

- ▶ Data $D = y_1, \dots, y_T$ with $\mathbb{E}[y_t] = \mu$.
- ▶ Goal: estimate μ with some estimator f_D to minimise
- ▶ MSE: $\mathbb{E}[(y - f_D)^2]$, the expected square difference between new samples and our guess.

Optimal estimate

To minimise the MSE, we use $f^* = \mu$. This gives us two ideas:

Empirical mean estimator:

- ▶ $f_D = \sum_{t=1}^T x_t / T$.
- ▶ $\mathbb{V}(f_D) = \mathbb{E}[f_D - \mu]^2 = 1/\sqrt{T}$
- ▶ $\mathbb{B}(f_D) = 0$.

Laplace mean estimator:

- ▶ $f_D = \sum_{t=1}^T (\lambda + x_t) / T$.
- ▶ $\mathbb{V}(f_D) = \mathbb{E}[f_D - \mu]^2 = \frac{1}{1+\sqrt{T}}$
- ▶ $\mathbb{B}(f_D) = O(1/T)$.

A proof of the bias/variance trade-off

- ▶ RV's $y_t \sim P$, $\mathbb{E}[y_t] = \mu$, $y_t = \mu + \epsilon_t$.
- ▶ Estimator f_D , $D = y_1, \dots, y_{t-1}$.

$$\begin{aligned}\mathbb{E}[(f_D - y_t)^2] &= \mathbb{E}[f_D^2] - 2\mathbb{E}[f_D y_t] + \mathbb{E}[y_t^2] \\ &= \mathbb{V}[f_D] + \mathbb{E}[f_D]^2 - 2\mathbb{E}[f_D y_t] + \mathbb{E}[y_t^2] \\ &= \mathbb{V}[f_D] + \mathbb{E}[f_D]^2 - 2\mathbb{E}[f_D]\mathbb{E}[y_t] + \mathbb{E}[y_t^2] \\ &= \mathbb{V}[f_D] + \mathbb{E}[f_D]^2 - 2\mathbb{E}[f_D]\mu + \mathbb{E}[y_t^2] \\ &= \mathbb{V}[f_D] + \mathbb{E}[f_D]^2 - 2\mathbb{E}[f_D]\mu + \mathbb{E}[(\mu + \epsilon_t)^2] \\ &= \mathbb{V}[f_D] + \mathbb{E}[f_D]^2 - 2\mathbb{E}[f_D]\mu + \mathbb{E}[\mu^2 + 2\mu\epsilon_t + \epsilon_t^2] \\ &= \mathbb{V}[f_D] + \mathbb{E}[f_D]^2 - 2\mathbb{E}[f_D]\mu + \mu^2 + \sigma^2 \\ &= \mathbb{V}[f_D] + (\mathbb{E}[f_D] - \mu)^2 + \sigma^2 \\ &= \mathbb{V}(f_D) + \mathbb{B}(f_D)^2 + \sigma^2\end{aligned}$$

Validation sets

Cross-validation

Bootstrapping

The wrong way to do XV for subset selection

1. Screen the predictors: find a subset of “good” predictors that show fairly strong (univariate) correlation with the class labels
1. Using just this subset of predictors, build a multivariate classifier.
2. Use cross-validation to estimate the unknown tuning parameters and to estimate the prediction error of the final model. Is this a correct application of cross-validation? Consider a scenario with $N = 50$ samples in two equal-sized classes, and $p = 5000$ quantitative predictors (standard Gaussian) that are independent of the class labels. The true (test) error rate of any classifier is 50%.

The right way to do XV for feature selection

1. Divide the samples into K cross-validation folds (groups) at random.
2. For each fold $k = 1, 2, \dots, K$ (a) Find a subset of “good” predictors that show fairly strong (univariate) correlation with the class labels, using all of the samples except those in fold k . (b) Using just this subset of predictors, build a multivariate classifier, using all of the samples except those in fold k . (c) Use the classifier to predict the class labels for the samples in fold k .

The perceptron algorithm

Input

- ▶ Feature space $X \subset \mathbb{R}^n$.
- ▶ Label space $Y = \{-1, 1\}$.
- ▶ Data (x_t, y_t) , $t \in [T]$, with $x_t \in X, y_t \in Y$.

Algorithm

- ▶ $w_1 = w_0$.
- ▶ For $t = 1, \dots, T$.
 - $a_t = \text{sgn}(w_t^\top x_t)$. – If $a_t \neq y_t$ — $w_{t+1} = w_t + y_t x_t$ – Else — $w_{t+1} = w_t$
- ▶ Return w_{T+1}

Theorem

The number of mistakes made by the perceptron algorithm is bounded by $(r/\rho)^2$, where $\|x_t\| \leq r$, $\rho \leq y_t(v^\top x_t)/\|v\|$ for some **margin** ρ and **hyperplane** v .

Perceptron examples

Example 1: One-dimensional data

- ▶ Done on the board
- ▶ Shows how the algorithm works.
- ▶ Demonstrates the idea of a margin

Example 2: Two-dimensional data

- ▶ See in-class programming exercise

Python concepts

Numpy

- ▶ `np.random.multivariate_normal()`: generate samples from an n-D normal distribution
- ▶ `np.random.choice()`: generate samples from a discrete distribution
- ▶ `np.zeros()`: generate an array of zeros
- ▶ `np.array()`: create an array from a list
- ▶ `np.block()`: make an array from nested lists
- ▶ `np.dot()`: calculate the dot (aka inner) product

matplotlib.pyplot

- ▶ `plt.plot()`: Plot lines and points
- ▶ `plt.axis()`: manipulate axes
- ▶ `plt.grid()`: show a grid
- ▶ `plt.show()`: display the plot

Gradient methods example

Estimate the expected value

$x_t \sim P$ with $\mathbb{E}_P[x_t] = \mu$.

Objective

$$\min_{\theta} \mathbb{E}_P[(x_t - \theta)^2].$$

Derivative

Idea: at the minimum the derivative should be zero.

$$d/d\theta \mathbb{E}_P[(x_t - \theta)^2] = \mathbb{E}_P[d/d\theta (x_t - \theta)^2] = \mathbb{E}_P[-2(x_t - \theta)] = \mathbb{E}_P[x_t] - 2\theta.$$

Setting the derivative to 0, we have $\theta = \mathbb{E}_P[x_t]$. This is a simple solution.

Real-world setting

- ▶ The objective function does not result in a simple solution
- ▶ The distribution P is not known.
- ▶ We can sample $x \sim P$.

Stochastic gradient for mean estimation

$$\begin{aligned}\frac{d}{d\theta} \mathbb{E}_P[(x - \theta)^2] &= \int_{-\infty}^{\infty} dP(x) \frac{d}{d\theta} (x - \theta)^2 \\ &= \frac{d}{d\theta} \int_{-\infty}^{\infty} dP(x) (x - \theta)^2\end{aligned}$$

Simple linear regression

Input and output

- ▶ Data pairs (x_t, y_t) , $t = 1, \dots, T$.
- ▶ Input $x_t \in \mathbb{R}^n$
- ▶ Output $y_t \in \mathbb{R}$.

Predicting the conditional mean $\mathbb{E}[y_t|x_t]$

- ▶ Parameters $\theta \in \mathbb{R}^n$
- ▶ Function $f_\theta : \mathbb{R}^n \rightarrow \mathbb{R}$, defined as

$$f_\theta(x_t) = \theta^\top x_t = \sum_{i=1}^n \theta_i x_{t,i}$$

Optimisation goal: Minimise mean-squared error.

$$\min_{\theta} \sum_{t=1}^T [y_t - \pi_{\theta}(x_t)]^2$$

How can we solve this problem?

Gradient descent algorithm

Minimising a function

$$\min_{\theta} f(\theta) \geq f(\theta') \forall \theta', \quad \theta^* = \arg \min_{\theta} f(\theta) \Rightarrow f(\theta^*) = \min_{\theta} f(\theta)$$

Gradient descent for minimisation

- ▶ Input θ_0
- ▶ For $n = 0, \dots, N$:
- ▶ $\theta_{n+1} = \theta_n - \eta_n \nabla_{\theta} f(\theta_n)$

Step-size η_n

- ▶ η_n fixed: for online learning
- ▶ $\eta_n = c/[c + n]$ for asymptotic convergence
- ▶ $\eta_n = \arg \min_{\eta} f(\theta_n + \eta \nabla_{\theta})$: Line search.

Gradient descent for squared error

Cost function

$$\ell(\theta) = \sum_{t=1}^T [y_t - \pi_{\theta}(x_t)]^2$$

Cost gradient

Using the chain rule of differentiation:

$$\begin{aligned}\nabla_{\theta} \ell(\theta) &= \nabla \sum_{t=1}^T [y_t - \pi_{\theta}(x_t)]^2 \\ &= \sum_{t=1}^T \nabla [y_t - \pi_{\theta}(x_t)]^2 \\ &= \sum_{t=1}^T 2[y_t - \pi_{\theta}(x_t)][-\nabla \pi_{\theta}(x_t)]^2\end{aligned}$$

Parameter gradient

Analytical Least-Squares Solution

Stochastic gradient descent algorithm

When f is an expectation

$$f(\theta) = \int_{\mathcal{X}} dP(x) g(x, \theta).$$

Replacing the expectation with a sample:

$$\begin{aligned} \nabla f(\theta) &= \int_{\mathcal{X}} dP(x) \nabla g(x, \theta) \\ &\approx \frac{1}{K} \sum_{k=1}^K \nabla g(x^{(k)}, \theta), \end{aligned} \quad x^{(k)} \sim P.$$

Layering and features

Fixed layers

- ▶ Fixed number of units, architecture and parameters
- ▶ Example 1: Feature transformation
- ▶ Example 2: Softmax layer

Adaptive layers

- ▶ Fixed units and architecture
- ▶ Adaptive parameters
- ▶ Example 1: Linear layer
- ▶ Example 2: Convolutional layers (e.g. images)

Softmax layer

- ▶ Features x
- ▶ Linear activation layer

$$z_i = \theta_i^\top x$$

Softmax output layer

We want to translate the real-valued z_i into probabilities:

$$y_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}.$$

Random projections

- ▶ Features x
- ▶ Hidden layer activation z
- ▶ Output y

Hidden layer: Random projection

Here we project the input into a high-dimensional space

$$z_i = \text{sgn}(\theta_i^\top x),$$

where $\Theta = [\theta_i]_{i=1}^m$.

The reason for random projections

- ▶ The high dimension makes it easier to learn.
- ▶ The randomness ensures we are not learning something spurious.

Back-propagation

The chain rule

$$f : X \rightarrow Z, \quad g : Z \rightarrow Y, \quad \frac{dg}{dx} = \frac{dg}{df} \frac{df}{dx}$$

Parametrised functions

$$f : \mathcal{W} \times X \rightarrow Z, \quad g : \Omega \times Z \rightarrow Y, \quad \pi = fg$$

(network mappings)

$$\ell(D, \pi) = \sum_{(x,y) \in D} [y - \pi(x)]^2 \quad (1)$$

Gradient descent with *back-propagation*

Apply the chain rule

$$\nabla_{w,\omega} \pi = \nabla_{\omega}$$

Neural architectures

Layers

- ▶ Input to layer $x \in R^n$
- ▶ Output from layer $z \in R^m$.

Linear layer

Transform the output of previous layers or features into either:

- ▶ A higher-dimensional space.
- ▶ A lower-dimensional space.
- ▶ They have adaptive parameters.
- ▶ Parameters can be dependent on each other for invariance (cf. convolution)

Non-linear layers

- ▶ Simple transformations of previous output
- ▶ Examples: Sigmoid, Softmax

Linear layer

Definition

This is a linear combination of inputs $x \in \mathbb{R}^n$ and parameter matrix

$$\mathbf{W} \in \mathbb{R}^{m \times n} \text{ where } \mathbf{W} = \begin{bmatrix} \mathbf{w}_1 \\ \vdots \\ \mathbf{w}_i \\ \vdots \\ \mathbf{w}_m \end{bmatrix} = \begin{bmatrix} w_{1,1} & \cdots & w_{1,j} & \cdots & w_{1,m} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ w_{i,1} & \cdots & w_{i,j} & \cdots & w_{i,m} \\ \vdots & \ddots & \vdots & \ddots & \vdots \\ w_{n,1} & \cdots & w_{i,j} & \cdots & w_{n,m} \end{bmatrix}$$

$$f(\mathbf{W}, \mathbf{x}) = \mathbf{W}\mathbf{x} \quad f_i(\mathbf{W}, \mathbf{x}) = \mathbf{w}_i \cdot \mathbf{x} = \sum_{j=1}^n w_{i,j} x_j,$$

Gradient

Each partial derivative is simple:

$$\frac{\partial}{\partial w_{i,j}} f_k(\mathbf{W}, \mathbf{x}) = x_j \mathbb{I}\{j = k\}$$

Sigmoid layer

Definition

This layer transforms each input non-linearly

$$f_j(\mathbf{x}) = 1/[1 + \exp(-x_j)] =$$

without looking at the other inputs.

Derivative

So let us ignore the other inputs for simplicity:

$$\frac{d}{dx} f(x) = \exp(-x)/[1 + \exp(-x)]^2$$

Softmax

The maximum margin classifier

Soft margins

Kernel methods

Probabilistic modelling

The problem

- ▶ Model family $\{P_\theta : \theta \in \Theta\}$
- ▶ Each model assigns a probability $P_\theta(x)$ to the data x .
- ▶ How can we estimate θ from x ?

Maximum Likelihood (ML) Estimation

$$\hat{\theta}(x) = \arg \max_{\theta} P_\theta(x).$$

Maximum A Posteriori (MAP) Estimation

Here we also need a prior distribution, but still estimate a single parameter:

- ▶ Prior $\beta(\theta)$, a distribution on Θ .
- ▶ $\hat{\theta}(x) = \arg \max_{\theta} P_\theta(x)\beta(\theta)$.

Bayesian Estimation

Here we estimate the complete distribution over parameters

- ▶ $\beta(\theta|x) = P_\theta(x)\beta(\theta) / \sum_{\theta'} P_{\theta'}(x)\beta(\theta')$

The Bernoulli distribution: Modelling a coin

Definition

If $x_t \sim \text{Bernoulli}(\theta)$ then $x_t = 1$ w.p. θ and $x_t = 0$ w.p. $1 - \theta$.

Likelihood function

$$P(x_1, \dots, x_T | \theta) = \prod_{t=1}^T P(x_t | \theta) = \prod_{t=1}^T \theta^{x_t} (1 - \theta)^{1 - x_t}$$

Maximum Likelihood Estimate

$$\arg \max_{\theta} P(x | \theta) = \arg \max_{\theta} \ln P(x | \theta).$$

$$\begin{aligned} \frac{d}{d\theta} \ln P(x | \theta) &= \frac{d}{d\theta} \left[\sum_t \ln P(x_t | \theta) \right] = \frac{d}{d\theta} \left[\sum_t \ln \theta^{x_t} (1 - \theta)^{1 - x_t} \right] \\ &= \frac{d}{d\theta} \left[\ln(\theta) \sum_t x_t + \ln(1 - \theta) \sum_t (1 - x_t) \right] \\ &= \frac{1}{\theta} \sum_t x_t - \frac{1}{1 - \theta} \sum_t (1 - x_t) \end{aligned}$$

Setting the derivative to zero:

$$\hat{\theta}_T = \frac{1}{T} \sum_{t=1}^T x_t$$

Bayesian Estimate

The prior distribution $P(\theta)$

$$\theta \sim \text{Beta}(\alpha_1, \alpha_0)$$

The likelihood function $P(\mathbf{x}|\theta)$

$$P(x_1, \dots, x_T | \theta) = \prod_{t=1}^T P(x_t | \theta)$$

The posterior distribution $P(\theta|\mathbf{x})$

$$\theta \sim \text{Beta}(\alpha_1 + \sum_{t=1}^T x_t, \alpha_0 + \sum_{t=1}^T x_t).$$

The Gaussian distribution: Modelling gambling gains

Discriminative modelling: general idea

- ▶ Data (x, y)
- ▶ Easier to model $P(y|x)$
- ▶ No need to model $P(x)$.

Examples

- ▶ Linear regression
- ▶ Logistic regression
- ▶ Multi-layer perceptron

Linear regression

Model

- ▶ $z = \theta^\top x$
- ▶ $p_\theta(y|x) = \frac{1}{\sqrt{2\pi}\sigma} \exp(-\frac{1}{2\sigma^2}|z - y|^2)$

Two-class classification: logistic regression

Model

- ▶ $z = \theta^\top x$
- ▶ $P_\theta(y = 1|x) = \frac{1}{1+e^{-z}}$

General

The problems of Machine Learning (1 week)

Learning as Optimisation (4 weeks)

Learning as Probabilistic Inference (4 weeks)

- Probabilistic Models

- Classification: Discriminative modelling

- Classification: Generative modelling

- Density estimation

Sequence modelling (2 weeks)

Reinforcement Learning (2 weeks)

Generative modelling

general idea

- ▶ Data (x, y) .
- ▶ Need to model $P(y|x)$.
- ▶ Model the complete data distribution: $P(x|y)$, $P(x)$, $P(y)$.
- ▶ Calculate $P(y|x) = \frac{P(x|y)P(y)}{P(x)}$.

Examples

- ▶ Naive Bayes classifier
- ▶ Gaussian Mixture Classifier

Modelling the data distribution

- ▶ Need to estimate the density $P(x|y)$ for each class y .

Classification: Naive Bayes Classifier

- ▶ Data (x, y)
- ▶ $x \in X$
- ▶ $y \in Y \subset \mathbb{N}$, N_i : amount of data from class i

Separately model each class

- ▶ Assume each class data comes from a different normal distribution
- ▶ $x|y = i \sim \text{Normal}(\mu_i, \sigma_i I)$
- ▶ For each class, calculate
 - ▶ Empirical mean $\hat{\mu}_i = \sum_{t: y_t = i} x_t / N_i$
 - ▶ Empirical variance $\hat{\sigma}_i$.

Decision rule

Use Bayes's theorem:

$$P(y|x) = P(x|y)P(y)/P(x),$$

choosing the y with largest posterior $P(y|x)$.

- ▶ $P(x|y = i) \propto \exp(-\|\hat{\mu}_i - x\|^2 / \hat{\sigma}_i^2)$

General idea

Parametric models

- ▶ Fixed histograms
- ▶ Gaussian Mixtures

Non-parametric models

- ▶ Variable-bin histograms
- ▶ Infinite Gaussian Mixture Model
- ▶ Kernel methods

Histograms

Fixed histogram

- ▶ Hyper-Parameters: number of bins
- ▶ Parameters: Number of points in each bin.

Variable histogram

- ▶ Hyper-parameters: Rule for constructing bins
- ▶ Generally \sqrt{n} points in each bin.

Gaussian Mixture Model

Hyperparameters:

- ▶ Number of Gaussian k .

Parameters:

- ▶ Multinomial distribution θ over Gaussians
- ▶ For each Gaussian i , center μ_i , covariance matrix Σ_i .

Model. For each point x_t :

- ▶ $c_t = i$ w.p. θ_i
- ▶ $x_t | c_t = i \sim \text{Normal}(\mu_i, \Sigma_i)$.

Algorithms:

- ▶ Expectation Maximisation
- ▶ Gradient Ascent
- ▶ Variational Bayesian Inference (with appropriate prior)

GMM with EM

Objective function: log-likelihood

$$\ln P(x|\theta, \mu, \Sigma) = \ln \sum_i \theta_i P(x|\mu_i, \sigma_i)$$

Expectation Step

Maximization Step

GMM Classifier

Base class: sklearn GaussianMixtureModel

- ▶ *fit()* only works for Density Estimation
- ▶ *predict()* only predicts cluster labels

Problem

- ▶ Create a GMMClassifier class
- ▶ *fit()* should take X, y, arguments
- ▶ *predict()* should predict class labels
- ▶ Hint: Use *predict_proba()* and multiple GMM models

The problem of sequence prediction

- ▶ Data x_1, x_2, x_3, \dots
- ▶ At time t , make a prediction a_t for x_t .

Auto-regressive models

General idea

- Predict x_t from the last k inputs

$$x_t \approx g(x_{t-k}, \dots, x_{t-1})$$

Optimisation view

We wish to minimise the difference between our predictions a_t and the next symbol

$$\sum_t (a_t - x_t)^2$$

Probabilistic view

We wish to model

$$P(x_t | x_{t-k}, \dots, x_{t-1})$$

Linear auto-regression

Simple time-series data

- ▶ Observations $x_t \in \mathbb{R}$
- ▶ Parameters $\boldsymbol{\theta} \in \mathbb{R}^k$

$$\hat{x}_t = \sum_i \theta_i x_{t-i}.$$

Multi-dimensional time-series data

- ▶ Observations $x_t \in \mathbb{R}^n$
- ▶ Parameters $\boldsymbol{\Theta} \in \mathbb{R}^{k \times n}$

$$\hat{x}_t = \sum_i \theta_i^\top x_{t-i} = \sum_{i,j} \theta_{i,j} x_{t-i}.$$

Recursive models

General idea

- ▶ Maintain an *internal state* z_t , which summarises what has been seen.

$$z_t = f(z_{t-1}, x_{t-1}) \quad (\text{change state})$$

- ▶ Make predictions using the internal state

$$\hat{x}_t = g(z_t) \quad (\text{predict})$$

Examples

- ▶ Hidden Markov models
- ▶ Recurrent Neural Networks

Hidden Markov Models: General setting

Variables

- ▶ State z_t
- ▶ Observations x_t

Parameters

- ▶ Transition θ
- ▶ Observation ψ

Distributions

- ▶ Transition distribution $P_{\theta}(z_{t+1}|z_t)$
- ▶ Observation distribution $P_{\psi}(x_t|z_t)$.

HMMs: Discrete case

Variables

- ▶ State $z_t \in [n]$
- ▶ Observation $x_t \in [m]$

Transition distribution

Multinomial with

$$P_{\theta}(z_{t+1} = j | z_t = i) = \theta_{i,j}$$

Observation distribution

Multinomial with

$$P_{\theta}(x_t = j | z_t = i) = \psi_{i,j}$$

HMMs: Continuous case

Variables

- ▶ State $z_t \in [n]$
- ▶ Observation $x_t \in \mathbb{R}^m$

Transition distribution

Multinomial with

$$P_{\theta}(z_{t+1} = j | z_t = i) = \theta_{i,j}$$

Observation distribution

Gaussian with

$$P_{\theta}(x_t = x | z_t = i) \propto \exp(-\|x - \psi_i\|)$$

Density Estimation with EM

HMM Estimation with EM

General

The problems of Machine Learning (1 week)

Learning as Optimisation (4 weeks)

Learning as Probabilistic Inference (4 weeks)

Sequence modelling (2 weeks)

Reinforcement Learning (2 weeks)